


# Introduction to Neural Networks

# Lecture 2 Recap

# Linear Regression

= a supervised learning method to find a linear model of the form

$$\hat{y}_i = \theta_0 + \sum_{j=1}^d x_{ij}\theta_j = \theta_0 + x_{i1}\theta_1 + x_{i2}\theta_2 + \dots + x_{id}\theta_d$$



Goal: find a model that explains a target y given the input x

# Logistic Regression

- Loss function

$$\mathcal{L}(y_i, \hat{y}_i) = -y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log[1 - \hat{y}_i]$$

~~\*~~ Cost function


$$\rightarrow \mathcal{C}(\theta) = - \sum_{i=1}^n (y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log[1 - \hat{y}_i])$$

Minimization

~~$\hat{y}_i = \sigma(x_i \theta)$~~   
predictions


*Understand*

# Linear vs Logistic Regression



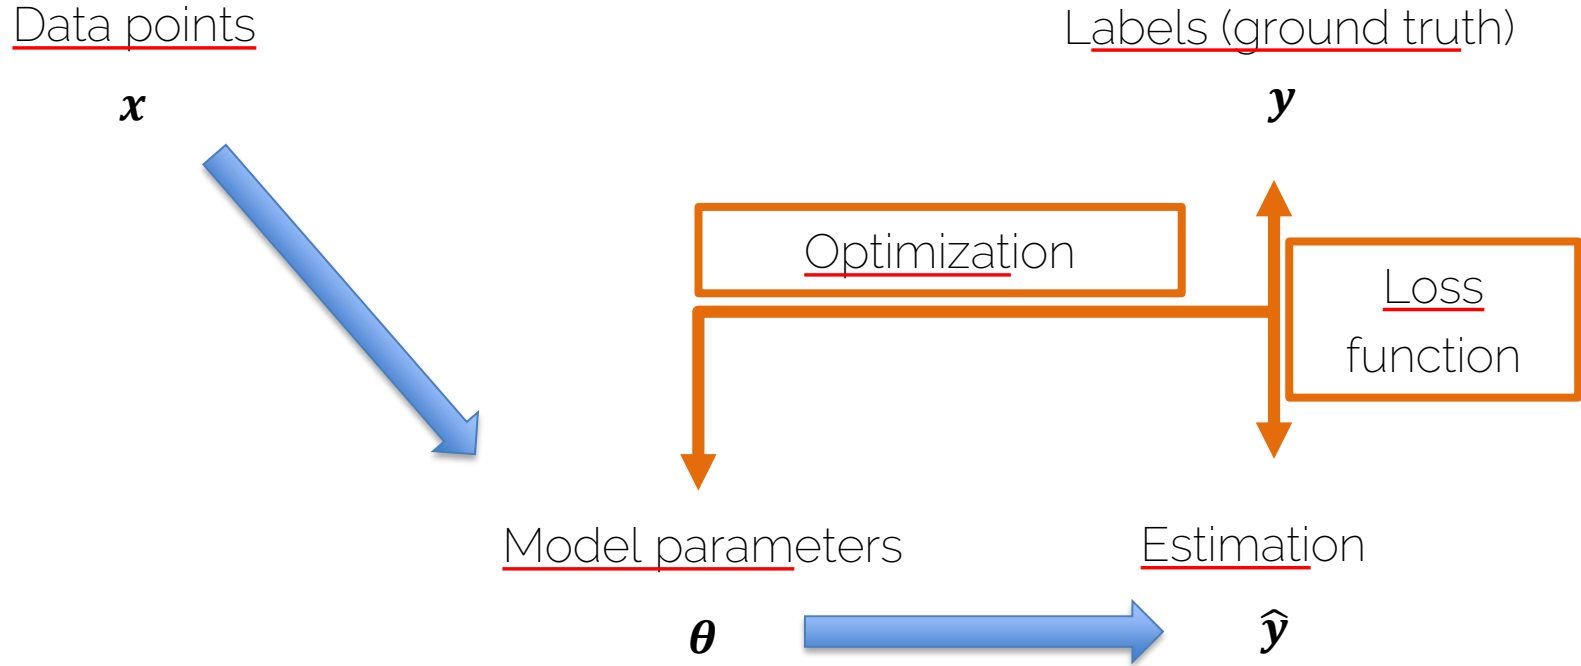
Predictions can exceed the range of the training samples

→ in the case of classification  $[0;1]$  this becomes a real issue



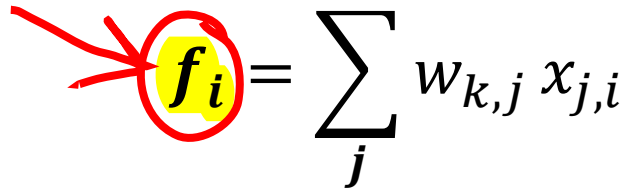
Predictions are guaranteed to be within  $[0;1]$

# How to obtain the Model?



# Linear Score Functions

- Linear score function as seen in linear regression


$$f_i = \sum_j w_{k,j} x_{j,i}$$

$$\mathbf{f} = \mathbf{W} \mathbf{x} \quad (\text{Matrix Notation})$$

\* Optimizing linear score is L2 optim

↳ optimal soln : mean


# Linear Score Functions on Images

- Linear score function  $f = \mathbf{W} \mathbf{x}$

*representation of the average image  
- so to say.*



On CIFAR-10



On ImageNet

Source: Li/Karpathy/Johnson


*object centered  
in the middle*

*\* Simple dataset  
might be described  
by linear fn.*




# Linear Score Functions?

Logistic Regression



Main  
Problem

Linear Separation Impossible!



not powerful enough to rep  
Complex fn

# Linear Score Functions?

- Can we make linear regression better?
  - Multiply with another weight matrix  $\mathbf{W}_2$

$$\hat{\mathbf{f}} = \mathbf{W}_2 \cdot \mathbf{f}$$

$$\hat{\mathbf{f}} = \mathbf{W}_2 \cdot \mathbf{W} \cdot \mathbf{x}$$

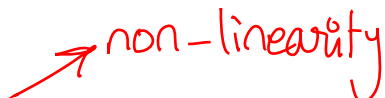
- Operation is still linear.

$$\hat{\mathbf{W}} = \mathbf{W}_2 \cdot \mathbf{W}$$

$$\hat{\mathbf{f}} = \hat{\mathbf{W}} \mathbf{x}$$


~~\*~~ Solution → add non-linearity!!

# Neural Network

- Linear score function  $f = Wx$
- Neural network is a nesting of 'functions'
  - 2-layers:  $f = W_2 \max(0, W_1 x)$   *non-linearity*
  - 3-layers:  $f = W_3 \max(0, W_2 \max(0, W_1 x))$
  - 4-layers:  $f = W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x)))$
  - 5-layers:  $f = W_5 \sigma(W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x))))$
  - ... up to hundreds of layers

# Introduction to Neural Networks


# History of Neural Networks



Source: [http://beamlab.org/deeplearning/2017/02/23/deep\\_learning\\_101\\_part1.html](http://beamlab.org/deeplearning/2017/02/23/deep_learning_101_part1.html)

# Neural Network

Logistic Regression



Neural Networks




# Neural Network

- Non-linear score function  $f = \dots (\max(0, W_1 x))$



On CIFAR-10




Visualizing activations of first layer.


Source: ConvNetJS

# Neural Network

1-layer network:  $f = Wx$




2-layer network:  $f = W_2 \max(0, W_1 x)$




Why is this structure useful?



# Neural Network




2-layer network:  $\mathbf{f} = \mathbf{W}_2 \max(\mathbf{0}, \mathbf{W}_1 \mathbf{x})$



dim


# Net of Artificial Neurons



$w_{l,m,n}$   
↙ ↘  
layer      neuron  
                 in layer  
                 ↘  
                 weight  
                 in neuron

NB Do not forget bias  
when counting weights

# Neural Network




Source: <https://towardsdatascience.com/training-deep-neural-networks-gfdb1964b964>


Important

# Activation Functions


Sigmoid:  $\sigma(x) = \frac{1}{1+e^{-x}}$




tanh:  $\tanh(x)$



ReLU:  $\max(0, x)$



Leaky ReLU:  $\max(0.1x, x)$



Parametric ReLU:  $\max(\alpha x, x)$

Maxout  $\max(w_1^T x + b_1, w_2^T x + b_2)$

→ ELU  $f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$

# Neural Network

$$f = W_3 \cdot (W_2 \cdot (W_1 \cdot x))$$

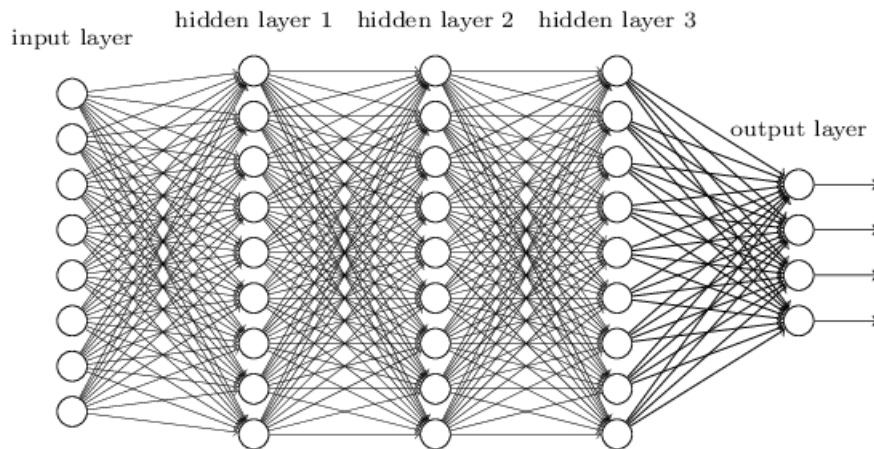
this is still  
linear

Why activation functions?


Simply concatenating linear layers would be so much cheaper...

# Neural Network

Why organize a neural network into layers?




# Biological Neurons




Credit: Stanford CS 231n

# Biological Neurons




*totally irrelevant*



Credit: Stanford CS 231n




# Artificial Neural Networks vs Brain



Artificial neural networks are **inspired** by the brain,  
but not even close in terms of complexity!

The comparison is great for the media and news articles however... 😊

# Artificial Neural Network



# Neural Network

- Summary  *Supervised*
  - Given a dataset with ground truth training pairs  $[x_i; y_i]$ ,

*idea*  Find optimal weights  $\mathbf{W}$  using stochastic gradient descent, such that the loss function is minimized

- *Compute gradients* with backpropagation (use batch-mode; more later)
- Iterate many times over training set (SGD; more later)


# Computational Graphs

# Computational Graphs

- Directional graph
- Matrix operations are represented as compute nodes.
- Vertex nodes are variables or operators like  $+$ ,  $-$ ,  $*$ ,  $/$ ,  $\log()$ ,  $\exp()$  ...
- Directional edges show flow of inputs to vertices

# Computational Graphs


- $f(x, y, z) = (x + y) \cdot z$



# Evaluation: Forward Pass

- $f(x, y, z) = (x + y) \cdot z$

Initialization  $x = 1, y = -3, z = 4$



# Computational Graphs

- Why discuss compute graphs?
- Neural networks have complicated architectures  
$$f = W_5 \sigma(W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x))))$$
- Lot of matrix operations!



Represent NN as computational graphs!

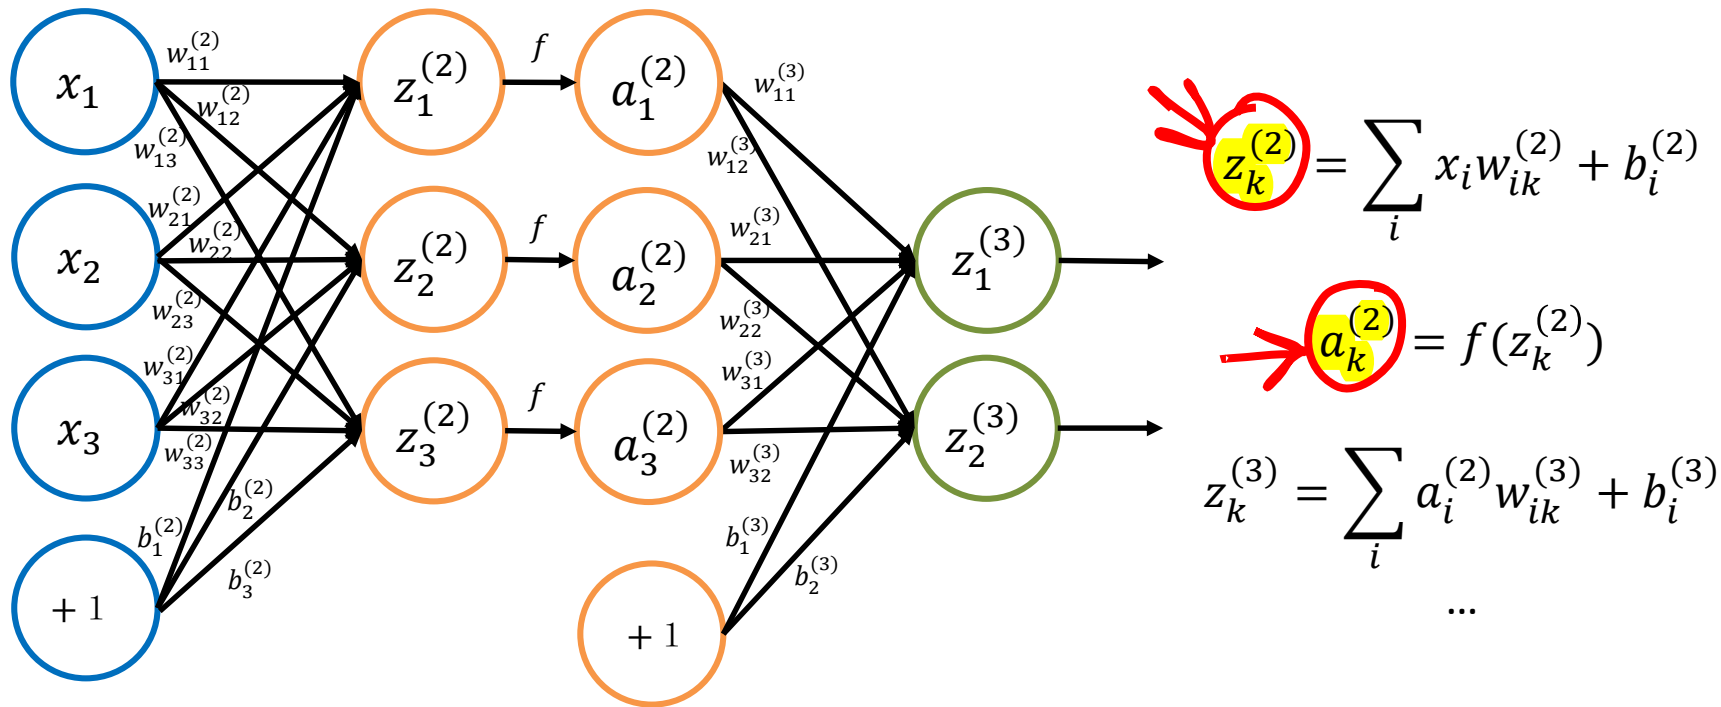


# Computational Graphs

A neural network can be represented as a computational graph...


- it has compute nodes (operations)
- it has edges that connect nodes (data flow)
- it is directional
- it can be organized into 'layers'

# Computational Graphs



# Computational Graphs

- From a set of neurons to a Structured Compute Pipeline




[Szegedy et al., CVPR'15] Going Deeper with Convolutions

Understand

# Computational Graphs

- The computation of Neural Network has further meanings:
  - The multiplication of  $W_i$  and  $x$ : encode input information
  - The activation function: select the key features




Source: <https://www.zybuluo.com/liuhui0803/note/981434>

# Computational Graphs

- The computations of Neural Networks have further meanings:

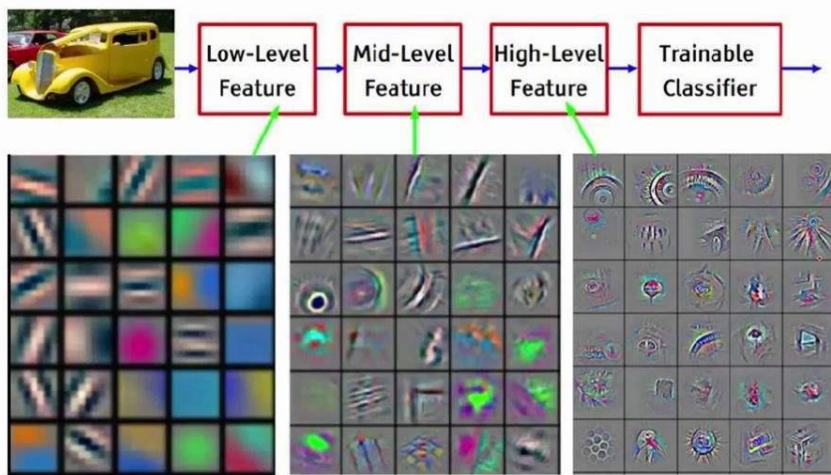
➔ The **convolutional layers**: extract useful features with shared weights



Source: <https://www.zcfy.cc/original/understanding-convolutions-colah-s-blog>

# Computational Graphs


- The computations of Neural Networks have further meanings:
  - The convolutional layers: extract useful features with shared weights



Source: <https://www.zybuluo.com/liuhui0803/note/981434>

# Loss Functions

# What's Next?




NB

We need a way to describe how close the network's outputs (= predictions) are to the targets!




# What's Next?

Idea: calculate a **'distance'** between prediction and target!



bad prediction



good prediction

# Loss Functions

- A function to measure the goodness of the predictions (or equivalently, the network's performance)

Intuitively, ...

- a large loss indicates bad predictions/performance (→ performance needs to be improved by training the model)
- ➔ the choice of the loss function depends on the concrete problem or the distribution of the target variable

# Regression Loss

- L1 Loss:

$$L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = \frac{1}{n} \sum_i^n \|y_i - \hat{y}_i\|_1$$

\* We normalize  
to avoid scaling  
issues

- MSE Loss:

L2 Loss

$$L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = \frac{1}{n} \sum_i^n \|y_i - \hat{y}_i\|_2^2$$

highly sensitive  
to outliers

# Binary Cross Entropy

- Loss function for binary (yes/no) classification

$$\rightarrow L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = - \sum_{i=1}^n (y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log[1 - \hat{y}_i])$$

① → training samples





Yes! (0.8)  
No! (0.2)

The network predicts the probability of the input belonging to the "yes" class!

# Cross Entropy

= loss function for multi-class classification

$$\rightarrow L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = - \sum_{i=1}^n \sum_{k=1}^k (y_{ik} \cdot \log \hat{y}_{ik})$$




This generalizes the binary case from the slide before!

$(K) \rightarrow \# \text{ classes}$


# More General Case

- Ground truth:  $\mathbf{y}$
- Prediction:  $\hat{\mathbf{y}}$
- Loss function:  $L(\mathbf{y}, \hat{\mathbf{y}})$
- Motivation: *it's all about optimization*
  - minimize the loss  $\Leftrightarrow$  find better predictions
  - predictions are generated by the NN
  - find better predictions  $\Leftrightarrow$  find better NN

# Initially




# During Training...






# During Training...




# Training Curve



Understand


# How to Find a Better NN?

Plotting loss curves against model parameters



# How to Find a Better NN?

- Loss function:  $L(\mathbf{y}, \hat{\mathbf{y}}) = L(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}))$
- Neural Network:  $f_{\boldsymbol{\theta}}(\mathbf{x})$
- Goal:
  - minimize the loss w. r. t.  $\boldsymbol{\theta}$




Optimization! We train compute graphs with some optimization techniques!

↳ for this lec, only gradient-based


# How to Find a Better NN?

- Minimize:  $L(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}))$  w.r.t.  $\boldsymbol{\theta}$
- In the context of NN, we use gradient-based optimization



# How to Find a Better NN?

- Minimize:  $L(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}))$  w.r.t.  $\boldsymbol{\theta}$



Learning rate

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}))$$

$$\boldsymbol{\theta}^* = \arg \min L(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}))$$

# How to Find a Better NN?

- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
- Given one layer NN with no activation function

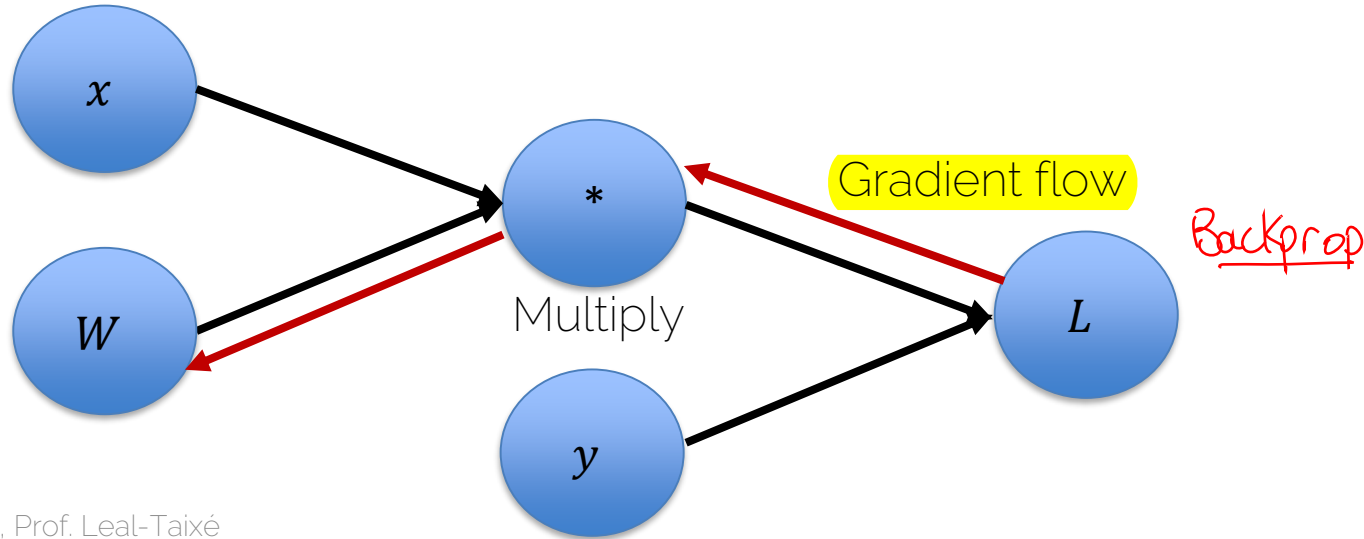
$$\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{W}\mathbf{x}, \quad \boldsymbol{\theta} = \mathbf{W}$$

Later  $\boldsymbol{\theta} = \{\mathbf{W}, \mathbf{b}\}$

- Given MSE Loss:  $L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = \frac{1}{n} \sum_i^n \|\mathbf{y}_i - \hat{\mathbf{y}}_i\|_2^2$

# How to Find a Better NN?

- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
- Given one layer NN with no activation function
- Given MSE Loss:  $L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = \frac{1}{n} \sum_i^n ||\mathbf{y}_i - \mathbf{W} \cdot \mathbf{x}_i||_2^2$






# How to Find a Better NN?

- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
- Given a one layer NN with no activation function

$$\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{W}\mathbf{x}, \quad \boldsymbol{\theta} = \mathbf{W}$$

- Given MSE Loss:  $L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = \frac{1}{n} \sum_i^n ||\mathbf{W} \cdot \mathbf{x}_i - \mathbf{y}_i||_2^2$

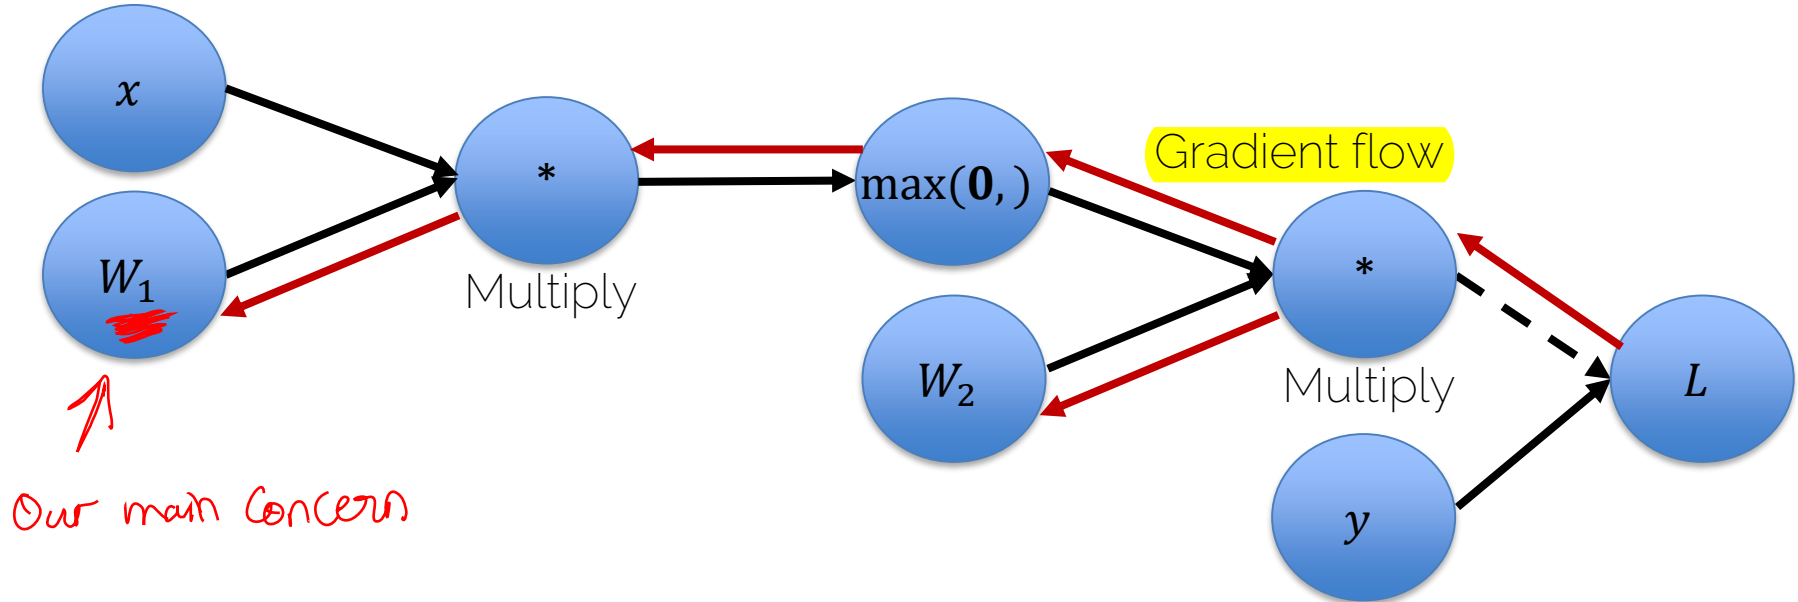

$$\nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x})) = \frac{1}{n} \sum_i^n (\mathbf{W} \cdot \mathbf{x}_i - \mathbf{y}_i) \cdot \mathbf{x}_i^T$$

# How to Find a Better NN?


- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
- Given a multi-layer NN with many activations
$$\mathbf{f} = \mathbf{W}_5 \sigma(\mathbf{W}_4 \tanh(\mathbf{W}_3, \max(\mathbf{0}, \mathbf{W}_2 \max(\mathbf{0}, \mathbf{W}_1 \mathbf{x}))))$$
- Gradient descent for  $L(\mathbf{y}, f_{\theta}(\mathbf{x}))$  w. r. t.  $\theta$ 
  - Need to propagate gradients from end to first layer ( $\mathbf{W}_1$ ).

# How to Find a Better NN?

- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
- Given multi-layer NN with many activations



# How to Find a Better NN?

- Given inputs  $\mathbf{x}$  and targets  $\mathbf{y}$
  - Given multilayer layer NN with many activations
$$\mathbf{f} = \mathbf{W}_5 \sigma(\mathbf{W}_4 \tanh(\mathbf{W}_3, \max(\mathbf{0}, \mathbf{W}_2 \max(\mathbf{0}, \mathbf{W}_1 \mathbf{x}))))$$
  - Gradient descent solution for  $L(\mathbf{y}, \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}))$  w. r. t.  $\boldsymbol{\theta}$ 
    - Need to propagate gradients from end to first layer ( $\mathbf{W}_1$ )
-  **Backpropagation:** Use chain rule to compute gradients
- Compute graphs come in handy!

# How to Find a Better NN?

- Why gradient descent?
  - Easy to compute using compute graphs
- Other methods include
  - Newtons method
  - L-BFGS
  - Adaptive moments
  - Conjugate gradient

– efficient with large networks

# Summary

- Neural Networks are computational graphs
- Goal: for a given train set, find optimal weights
- **Optimization** is done using gradient-based solvers
  - Many options (more in the next lectures)
- **Gradients** are computed via backpropagation
  - Nice because can easily modularize complex functions



# Next Lectures

- Next Lecture:
  - Backpropagation and optimization of Neural Networks
- Check for updates on website/moodle regarding exercises

See you next week 😊



# Further Reading

- Optimization:
  - <http://cs231n.github.io/optimization-1/>
  - <http://www.deeplearningbook.org/contents/optimization.html>
- General concepts:
  - Pattern Recognition and Machine Learning – C. Bishop
  - <http://www.deeplearningbook.org/>