



# Introduction to Deep Learning (I2DL)

## Exercise 9: Facial Keypoint Detection

# Overview

- **Optional Exercise: CIFAR-10**
  - Case study of two submitted solutions
- **Convolutional Layers**
  - Recap
  - Changes to Dropout & Batchnorm
- **Submission: Facial Keypoint Detection**
  - Deadline: January 27, 2021 15.59



# Case Study:

# Optional Exercise CIFAR-10

# Optional Exercise: Summary

- Image classifier on CIFAR-10 dataset
- CIFAR-10: Ten classes ('plane', 'car', 'bird', 'cat', ...)
- Pytorch Lightning
- Passing Criteria: 50%
- Restrictions in this exercise:

*\* The size of your final model must be less than 20 MB, which is approximately equivalent to 5 Mio. params. Note that this limit is quite lenient, you will probably need much less parameters!*

*Also, don't use convolutional layers as they've not been covered yet in the lecture and build your network with fully connected layers (`nn.Linear()`)!*

# Case Study: Leaderboard

## Leaderboard: Optional: CIFAR-10 (EX7)

Rank	User	Score	Pass
#1	w1251	68.91	✓
#2	w1465	60.00	✓
#3	w1567	59.48	✓
#4	w1710	57.48	✓
#5	w1258	56.62	✓
#6	w1232	56.05	✓
#7	w1829	55.80	✓
#8	w1571	55.43	✓
#9	w1328	55.19	✓
#10	w1337	54.84	✓

Submission Leaderboard Optional Exercise: CIFAR10 (17.01.2021)

# Case Study #1: Model

```
#####  
# TODO: Initialize your model! #  
#####  
  
self.model = nn.Sequential(  
    nn.Conv2d(3, 6, 3, padding=1),  
    torch.nn.BatchNorm2d(num_features=6),  
    nn.ReLU(),  
    nn.Conv2d(6, 9, 3, padding=1),  
    torch.nn.BatchNorm2d(num_features=9),  
    nn.ReLU(),  
    nn.Conv2d(9, 9, 5, padding=2),  
    torch.nn.BatchNorm2d(num_features=9),  
    nn.ReLU(),  
    nn.Conv2d(9, 9, 5, padding=2),  
    torch.nn.BatchNorm2d(num_features=9),  
    nn.ReLU(),  
    nn.Conv2d(9, 18, 5, stride=2, padding=2),  
    torch.nn.BatchNorm2d(num_features=18),  
    nn.ReLU(),  
    nn.Conv2d(18, 36, 5, stride=2, padding=2),  
    torch.nn.BatchNorm2d(num_features=36),  
    nn.ReLU(),  
    nn.Conv2d(36, 72, 3, stride=1, padding=1),  
    torch.nn.BatchNorm2d(num_features=72),  
    nn.ReLU(),  
    torch.nn.AvgPool2d((8,8)),  
    Lambda(lambda x: torch.squeeze(x)),  
    torch.nn.Linear(72, num_classes),  
    torch.nn.Softmax(dim=1)  
)
```

Also, don't use convolutional layers as they've not been covered yet in the lecture and build your network with fully connected layers (`nn.Linear()`)!

```
nn.Conv2d(18, 36, 5, stride=2, padding=2),  
torch.nn.BatchNorm2d(num_features=36),  
nn.ReLU(),  
nn.Conv2d(36, 72, 3, stride=1, padding=1),  
torch.nn.BatchNorm2d(num_features=72),  
nn.ReLU(),  
torch.nn.AvgPool2d((8,8)),  
Lambda(lambda x: torch.squeeze(x)),  
torch.nn.Linear(72, num_classes),  
torch.nn.Softmax(dim=1)
```

# Case Study #2: Model

```
#####  
# TODO: Initialize your model!                                     #  
#####  
  
modules = []  
for _ in range(hparams['num_layers']-2):  
    modules.append(nn.Linear(self.hparams['n_hidden'], self.hparams['n_hidden']))  
    modules.append(nn.ReLU())  
  
self.model = nn.Sequential(  
    nn.Linear(input_size, self.hparams['n_hidden']),  
    nn.ReLU(),  
    *modules,  
    nn.Linear(self.hparams['n_hidden'], num_classes),  
)  
  
#####  
#                               END OF YOUR CODE                               #  
#####
```

## Model Hyperparameters:

- Number of nodes in hidden layers
- Number of layers

## Very Simple Network Architecture

- No BatchNorm Layers
- No Dropout Layers
- Finally: 3 Blocks ( Linear (+ ReLu))

## Default initialization for Linear Layers

# Pytorch Default Weight Initialization

CLASS `torch.nn.Linear(in_features, out_features, bias=True)`

[SOURCE]

Applies a linear transformation to the incoming data:  $y = xA^T + b$

## Parameters

- **in\_features** – size of each input sample
- **out\_features** – size of each output sample
- **bias** – If set to `False`, the layer will not learn an additive bias. Default: `True`

## Shape:

- Input:  $(N, *, H_{in})$  where  $*$  means any number of additional dimensions and  $H_{in} = \text{in\_features}$
- Output:  $(N, *, H_{out})$  where all but the last dimension are the same shape as the input and  $H_{out} = \text{out\_features}$ .

## Variables

- **-Linear.weight** – the learnable weights of the module of shape  $(\text{out\_features}, \text{in\_features})$ . The values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ , where  $k = \frac{1}{\text{in\_features}}$
- **-Linear.bias** – the learnable bias of the module of shape  $(\text{out\_features})$ . If `bias` is `True`, the values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where  $k = \frac{1}{\text{in\_features}}$

Xavier/2 Init  
in comparison

$$\text{Var}(w_i) = \frac{2}{\text{fan\_in}}$$



# Case Study #2: Model

```
#####
# TODO: Initialize your model!                                     #
#####

modules = []
for _ in range(hparams['num_layers']-2):
    modules.append(nn.Linear(self.hparams['n_hidden'], self.hparams['n_hidden']))
    modules.append(nn.ReLU())

self.model = nn.Sequential(
    nn.Linear(input_size, self.hparams['n_hidden']),
    nn.ReLU(),
    *modules,
    nn.Linear(self.hparams['n_hidden'], num_classes),
)

#####
#                               END OF YOUR CODE                               #
#####
```

## Model Hyperparameters:

- Number of nodes in hidden layers
- Number of layers

## Default initialization for Linear Layers

- ReLu kills half of the data, so Xavier/2 initialization could be beneficial (see Lecture 7)
- Pytorch: `torch.nn.init.xavier_normal_` (takes `in_f` and `out_f` into consideration)

# Case Study #2: Transforms

```
my_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean, std),
    transforms.RandomHorizontalFlip(p=0.5),
    RandomTranslation(prob=0.5),
    transforms.RandomChoice([
        RandomSpeckle(std=0.2, prob=0.5),
        SaltandPepper(prob=0.5),
        #transforms.GaussianBlur(kernel_size=5, sigma=(0.1, 2.0)),
        #transforms.RandomRotation((90, 90), resample=False, expand=False, center=None, fill=None),
        #transforms.RandomRotation((-90, -90), resample=False, expand=False, center=None, fill=None),
        #transforms.RandomRotation((180, 180), resample=False, expand=False, center=None, fill=None),
        #transforms.RandomErasing(p=0.2, scale=(0.02, 0.2), ratio=(0.3, 3.3), value=0.05, inplace=False)
    ])
])
```

# Case Study #2: Hyperparameter Tuning

## Random Search

```
from exercise_code.MyPytorchModel import MyPytorchModel
from exercise_code.Util import printModelInfo
from exercise_code.Util import test_and_save
from pytorch_lightning.callbacks.early_stopping import EarlyStopping
import random
from math import log10

for i in range(100):
    early_stop_callback = EarlyStopping(
        monitor='val_loss',
        min_delta=0.0,
        patience=7,
        verbose=False,
        mode='min'
    )

    hparams = {}
    sample = random.uniform(log10(2.5e-4), log10(9e-4))
    lr = 10**(sample)
    num_layers = random.choice([2, 3, 4, 5])
    if num_layers == 2:
        hidden_end = 1600
    if num_layers == 3:
        hidden_end = 1170
    if num_layers == 4:
        hidden_end = 980
    if num_layers == 5:
        hidden_end = 870
    hparams = {
        'lr': lr,
        'decay': 0.0, #random.uniform(0.0, 0.3),
        'num_layers': 3, #num_layers,
        'n_hidden': 725, #random.choice([703, 725]), #random.randint(700,
        'batch_size': 2048, #random.choice([512, 1024, 2048]), #random.ch
        'num_workers': 3
    }
    print(hparams)
```

```
model = MyPytorchModel(hparams)
model.prepare_data()
_ = printModelInfo(model)

if hparams['batch_size'] == 512:
    epochs = 57
if hparams['batch_size'] == 1024:
    epochs = 50
if hparams['batch_size'] == 2048:
    epochs = 65

trainer = None
trainer = pl.Trainer(
    #callbacks=[early_stop_callback],
    precision=16,
    weights_summary=None,
    max_epochs=65, #epochs,
    #profiler='simple',
    progress_bar_refresh_rate=10,
    gpus=1
)
trainer.fit(model)
test_and_save(model)
```

# Case Study #2: Final Hyperparameters

```
hparams = {}  
sample = random.uniform(log10(2.5e-4), log10(9e-4))#random.uniform(log10(5e-6), log10(6e-3))  
lr = 10**(sample)
```


```
hparams = {  
    'lr': lr,  
    'decay': 0.0,#random.uniform(0.0, 0.3),  
    'num_layers': 3,#num_layers,  
    'n_hidden': 725,#random.choice([703, 725]),#random.randint(700,800),#random.randint(100, hidden_end),  
    'batch_size': 2048,#random.choice([512, 1024, 2048]),#random.choice([32, 64, 128, 256, 512, 1024, 2048]),  
    'num_workers': 3  
}
```

Take away: Always start with simple networks, you can already achieve quite good results

# Convolutional Layers



# Recap: Fully-Connected Layers

- **Regular Neural Networks:** Receive an input vector and transform it through a series of hidden layers.
- **Fully connected layers:** Each layer is made up of a set of neurons, where each single neuron is connected to all neurons in the previous layer



# Convolutional Layers

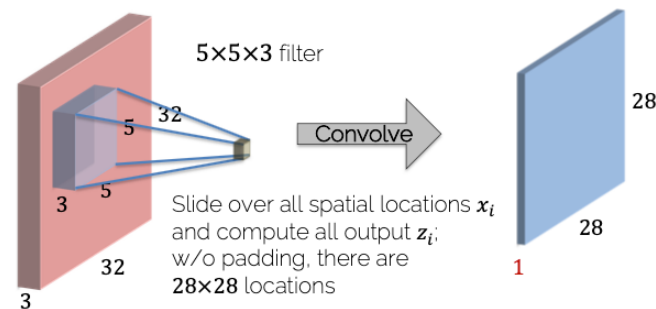
- **Assumption:** Input to our Network are images
- **Disadvantage:** Normal sized images are more likely to produce the right situation



Can we reduce the number of weights in our architecture?

# Convolutional Layers

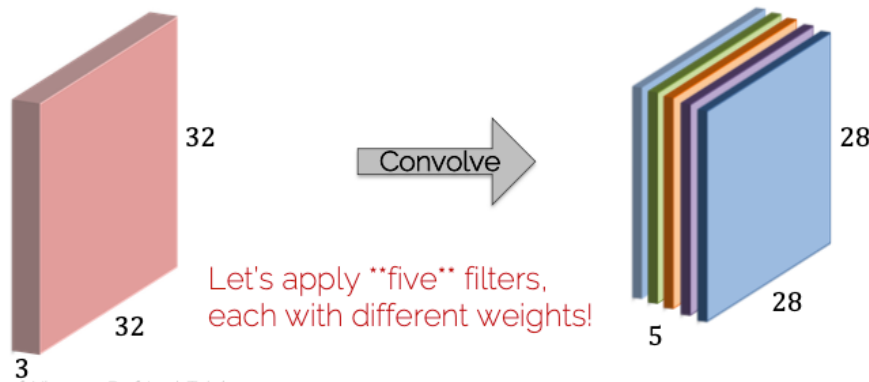
- **Assumption:** Input to our Network are images
- **Advantage:** We can analyze the image by looking at different region instead of looking at the whole image
- **Idea:** Sliding filter over the input image (convolution) instead of matrix multiplication






# Convolutional Layers

- **Assumption:** Input to our Network are images
- **Filters:** Sliding window with the same filter parameters to extract image features
  - Concept of weight sharing
  - Extract same features independent of location





# Fully Connected vs Convolution



# Convolutional Layers: BatchNorm and Dropout

# Fully Connected vs Convolution

- Output Fully-Connected layer: One layer of neurons, independent
- Output Convolutional Layer: Neurons arranged in 3 dimensions



# Recap: Batch Normalization

- Batchnorm for regular neural networks
  - Input size (N, D)
  - Compute minibatch mean and variance across N (i.e. we compute mean/var for each feature dimension)

**Input:**  $x : N \times D$

**Learnable params:**

$$\gamma, \beta : D$$

**Intermediates:**  $\mu, \sigma : D$   
 $\hat{x} : N \times D$

**Output:**  $y : N \times D$

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

# Recap: Batch Normalization

- Batchnorm for regular neural networks
  - Input size (N, D)
  - Compute minibatch mean and variance across N (i.e. we compute mean/var for each feature dimension)

Batch Normalization for  
**fully-connected** networks

**$\mathbf{x}$ :  $\mathbf{N} \times \mathbf{D}$**

Normalize



**$\boldsymbol{\mu}, \boldsymbol{\sigma}$ :  $1 \times \mathbf{D}$**

**$\boldsymbol{\gamma}, \boldsymbol{\beta}$ :  $1 \times \mathbf{D}$**


**$\mathbf{y} = \boldsymbol{\gamma}(\mathbf{x} - \boldsymbol{\mu}) / \boldsymbol{\sigma} + \boldsymbol{\beta}$**

# Spatial Batch Normalization

- Batchnorm for convolutional NN = spatial batchnorm
  - Input size (N, C, W, H)
  - Compute minibatch mean and variance across N, W, H (i.e. we compute mean/var for each channel C)

Batch Normalization for  
**convolutional** networks  
(Spatial Batchnorm, BatchNorm2D)


$$\begin{aligned} \mathbf{x} &: \mathbf{N} \times \mathbf{C} \times \mathbf{H} \times \mathbf{W} \\ \text{Normalize} \quad & \downarrow \quad \downarrow \quad \downarrow \\ \boldsymbol{\mu}, \boldsymbol{\sigma} &: \mathbf{1} \times \mathbf{C} \times \mathbf{1} \times \mathbf{1} \\ \boldsymbol{\gamma}, \boldsymbol{\beta} &: \mathbf{1} \times \mathbf{C} \times \mathbf{1} \times \mathbf{1} \\ \mathbf{y} &= \boldsymbol{\gamma}(\mathbf{x} - \boldsymbol{\mu}) / \boldsymbol{\sigma} + \boldsymbol{\beta} \end{aligned}$$




# Dropout for convolutional layers

- **Regular Dropout:** Deactivating specific neurons in the networks (one neuron “looks” at whole image)
- **Dropout Convolutional Layers:** Standard neuron-level dropout (i.e. randomly dropping a unit with a certain probability) does not improve performance in convolutional NN
- **Variant:** Spatial Dropout randomly sets entire feature maps to zero

Standard Dropout



Spatial Dropout






# Exercise 9:

# Facial Keypoints Detection

# Submission: Facial Keypoints



Input:  
(1, 96, 96) grayscale image

Output:  
(2, 15) keypoint coordinates

# Submission: Metric

```
def evaluate_model(model, dataset):  
    model.eval()  
    criterion = torch.nn.MSELoss()  
    dataloader = DataLoader(dataset, batch_size=1, shuffle=False)  
    loss = 0  
    for batch in dataloader:  
        image, keypoints = batch["image"], batch["keypoints"]  
        predicted_keypoints = model(image).view(-1,15,2)  
        loss += criterion(  
            torch.squeeze(keypoints),  
            torch.squeeze(predicted_keypoints)  
        ).item()  
    return 1.0 / (2 * (loss/len(dataloader)))  
  
print("Score:", evaluate_model(dummy_model, val_dataset))
```

# Submission: Details

- **Submission Start: January 21, 2021 13.00**
- **Submission Deadline : January 27, 2021 15.59**
- **Your model's evaluation score is all that counts!**
  - Evaluation score:  $1 / (2 * \text{MSE})$
  - **A score of at least 100** to pass the submission

Good luck &  
see you next week

