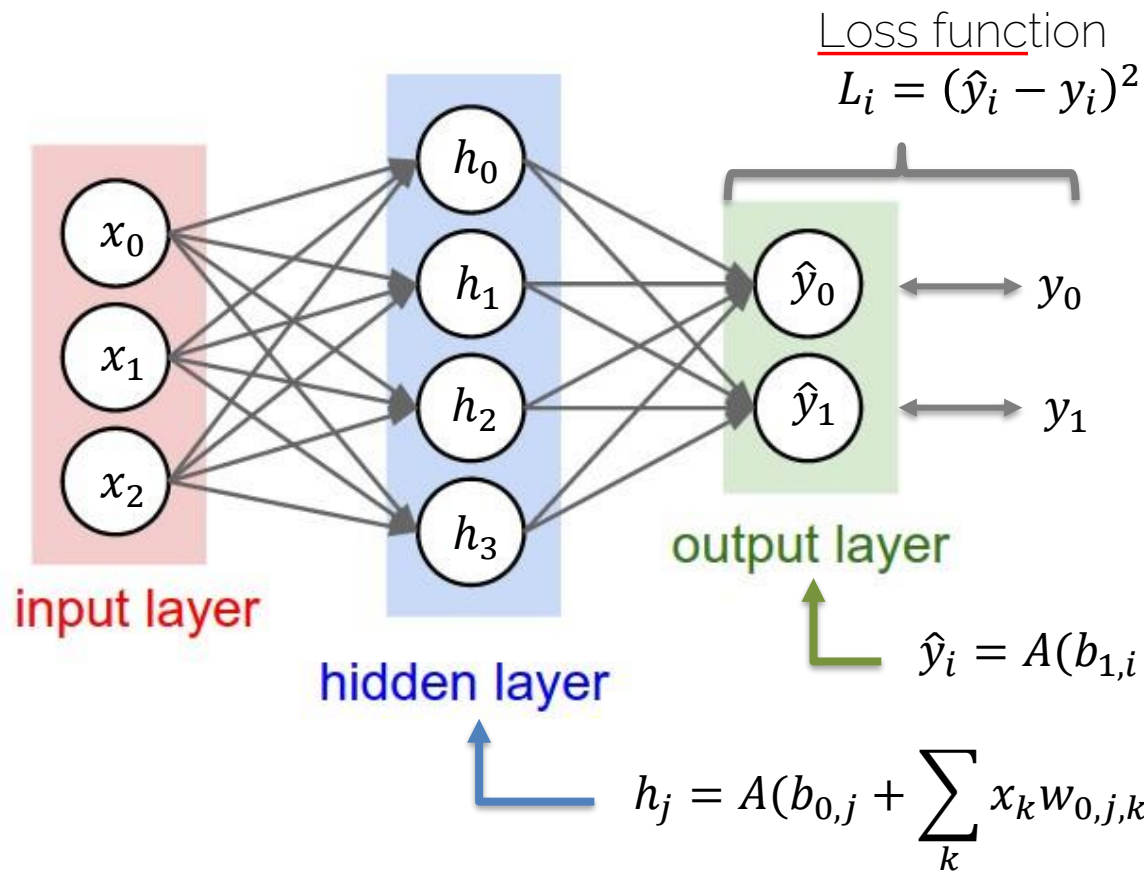


# Training Neural Networks

# Lecture 5 Recap

# Gradient Descent for Neural Networks



$$\nabla_{W,b} f_{\{x,y\}}(W) = \begin{bmatrix} \frac{\partial f}{\partial w_{0,0,0}} \\ \dots \\ \frac{\partial f}{\partial w_{l,m,n}} \\ \dots \\ \frac{\partial f}{\partial b_{l,m}} \end{bmatrix}$$

optimizing for weights & biases

Just simple:  
 $A(x) = \max(0, x)$

# Stochastic Gradient Descent (SGD)

~~\*~~  $\theta^{k+1} = \theta^k - \alpha \nabla_{\theta} L(\theta^k, \mathbf{x}_{\{1..m\}}, \mathbf{y}_{\{1..m\}})$

$k$  now refers to  $k$ -th iteration

$$\nabla_{\theta} L = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} L_i$$

$m$  training samples in the current minibatch

Gradient for the  $k$ -th minibatch

# Gradient Descent with Momentum

In case grads are  
going in the same dir,  
we are accumulating them  
i.e. going faster

- $\mathbf{v}^{k+1} = \beta \cdot \mathbf{v}^k + \nabla_{\theta} L(\theta^k)$

accumulation rate  
(‘friction’, momentum)

velocity

Gradient of current minibatch

- $\theta^{k+1} = \theta^k - \alpha \cdot \mathbf{v}^{k+1}$

model


learning rate

velocity

Exponentially-weighted average of gradient

Important: velocity  $\mathbf{v}^k$  is vector-valued!

# RMSProp



(Uncentered) variance of gradients  
→ second momentum

$$\mathbf{s}^{k+1} = \beta \cdot \mathbf{s}^k + (1 - \beta)[\nabla_{\theta} L \circ \nabla_{\theta} L]$$

We're dividing by square gradients:

- Division in Y-Direction will be large
- Division in X-Direction will be small

$$\boldsymbol{\theta}^{k+1} = \boldsymbol{\theta}^k - \alpha \cdot \frac{\nabla_{\theta} L}{\sqrt{\mathbf{s}^{k+1} + \epsilon}}$$

Can increase learning rate!

# Adam

- Combines Momentum and RMSProp

•  $\underset{\text{"mean"}}{\mathbf{m}^{k+1}} = \beta_1 \cdot \mathbf{m}^k + (1 - \beta_1) \nabla_{\theta} L(\theta^k)$      •  $\underset{\text{"Variance"}}{\mathbf{v}^{k+1}} = \beta_2 \cdot \mathbf{v}^k + (1 - \beta_2) [\nabla_{\theta} L(\theta^k) \circ \nabla_{\theta} L(\theta^k)]$

- $\mathbf{m}^{k+1}$  and  $\mathbf{v}^{k+1}$  are initialized with zero
  - bias towards zero
  - Typically, bias-corrected moment updates


$$\hat{\mathbf{m}}^{k+1} = \frac{\mathbf{m}^{k+1}}{1 - \beta_1^{k+1}} \quad \hat{\mathbf{v}}^{k+1} = \frac{\mathbf{v}^{k+1}}{1 - \beta_2^{k+1}} \quad \longrightarrow \quad \theta^{k+1} = \theta^k - \alpha \cdot \frac{\hat{\mathbf{m}}^{k+1}}{\sqrt{\hat{\mathbf{v}}^{k+1} + \epsilon}}$$

# Training Neural Nets



# Learning Rate: Implications

- What if too high?
- What if too low?




Source: <http://cs231n.github.io/neural-networks-3/>

# Learning Rate

*Ideally*

Need high learning rate when far away



Need low learning rate when close

# Learning Rate Decay

- $\alpha = \frac{1}{1 + \text{decay\_rate} \cdot \text{epoch}} \cdot \alpha_0$

*$\alpha$  gets smaller, as epoch increase*

– E.g.,  $\alpha_0 = 0.1$ , decay\_rate = 1.0

→ Epoch 0: 0.1


→ Epoch 1: 0.05

→ Epoch 2: 0.033

→ Epoch 3: 0.025

...

Learning Rate over Epochs



Same idea : start falling quickly "high  $\alpha$ ",  
& then relatively keep it constant

Many options:

- **Step decay**  $\alpha = \alpha - t \cdot \alpha$  (only every n steps)
  - T is decay rate (often 0.5)
- **Exponential decay**  $\alpha = t^{epoch} \cdot \alpha_0$ 
  - t is decay rate ( $t < 1.0$ )
- **$\alpha = \frac{t}{\sqrt{epoch}} \cdot \alpha_0$** 
  - t is decay rate
- Etc.

# Training Schedule

Manually specify learning rate for entire training process

- Manually set learning rate every n-epochs
- How?
  - Trial and error (the hard way)
  - Some experience (only generalizes to some degree)

Consider: #epochs, training set size, network size, etc.

# Basic Recipe for Training

- Given ground dataset with ground labels
  - $\{x_i, y_i\}$ 
    - $x_i$  is the  $i^{th}$  training image, with label  $y_i$
    - Often  $\dim(x) \gg \dim(y)$  (e.g., for classification)
    - $i$  is often in the 100-thousands or millions
  - Take network  $f$  and its parameters  $w, b$
  - ✗ Use SGD (or variation) to find optimal parameters  $w, b$ 
    - Gradients from backpropagation

# Gradient Descent on Train Set

- Given large train set with ( $n$ ) training samples  $\{\mathbf{x}_i, \mathbf{y}_i\}$ 
  - Let's say 1 million labeled images
  - Let's say our network has 500k parameters

- Gradient has 500k dimensions

- $n = 1 \text{ million}$

**\*** Extremely expensive to compute

*\* Even with minibatches, it is still expensive!*

*↳ majority of computation is gradient.*

# Learning

\* Learning means generalization to unknown dataset

- (So far no 'real' learning)
- I.e., train on known dataset  $\rightarrow$  test with optimized parameters on unknown dataset

- Basically, we hope that based on the train set, the optimized parameters will give similar results on different data (i.e., test data)

\* **GOAL**  $\rightarrow$  getting good descriptors, that are agnostic to specific dataset  
 $\downarrow$  making training harder i.e. avoiding data memorization



- Training set ('train'):
    - Use for training your neural network
  - Validation set ('val'):
    - Hyperparameter optimization
    - Check generalization progress
  - Test set ('test'):
    - Only for the very end
- ➔ NEVER TOUCH DURING DEVELOPMENT OR TRAINING

- Typical splits
  - Train (60%), Val (20%), Test (20%)
  - Train (80%), Val (10%), Test (10%)


Important

- \* During training:
  - Train error comes from average minibatch error
  - ➔ Typically take subset of validation every n iterations

\* Doing so is not expensive! Computing grad through backprop is by an order of magnitude more expensive.

# Basic Recipe for Machine Learning

- Split your data



- Split your data



Example scenario

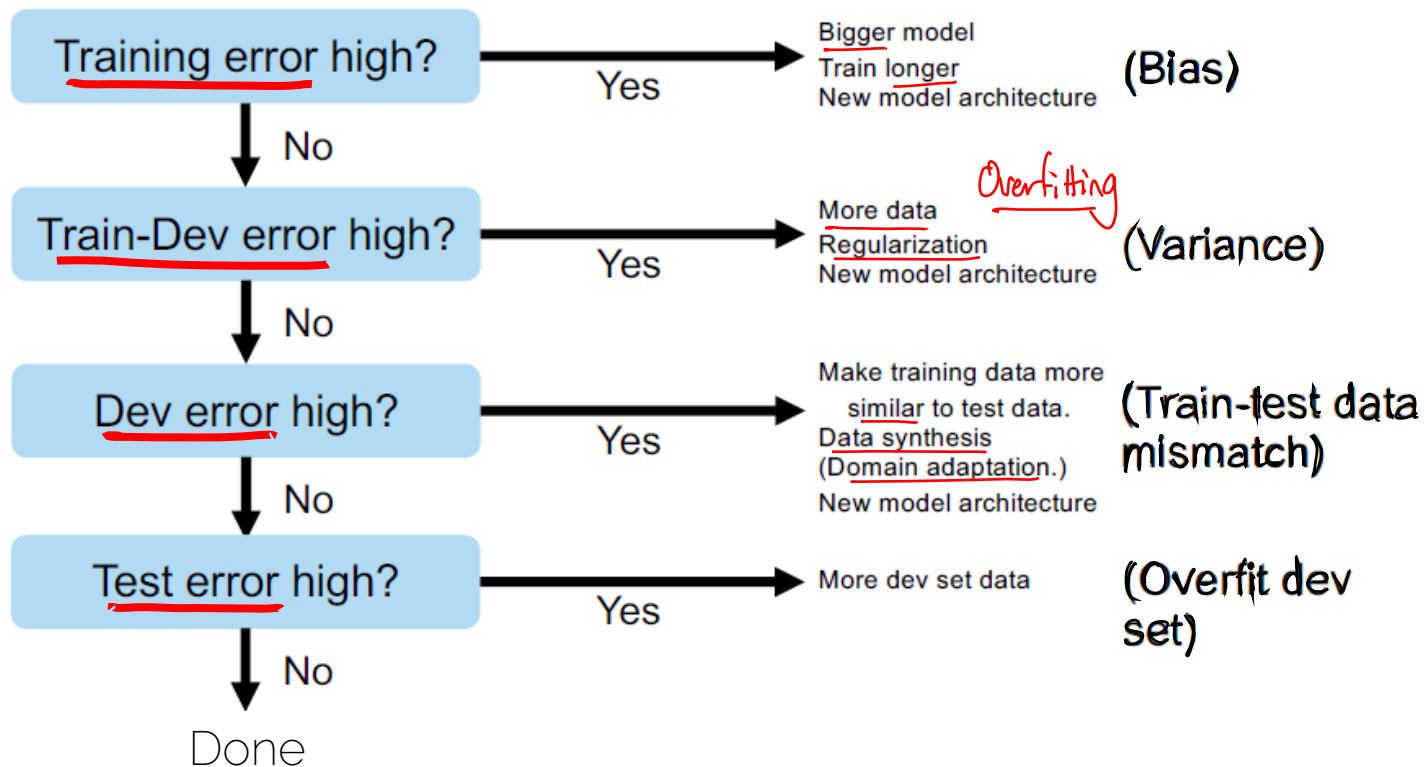
Ground truth error ..... 1%

Training set error ..... 5%

Val/test set error ..... 8%




Very Important!




Credits: A. Ng


# Over- and Underfitting



Underfitted




Appropriate



Overfitted

Source: Deep Learning by Adam Gibson, Josh Patterson, O'Reilly Media Inc., 2017




Source: <https://srdas.github.io/DLBook/ImprovingModelGeneralization.html>

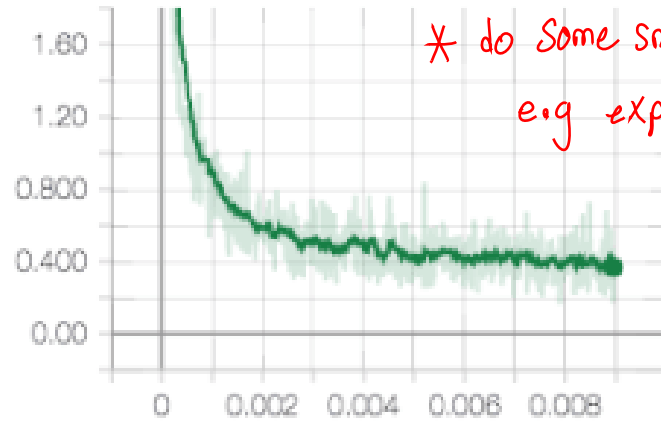
# Learning Curves

- Training graphs

- Accuracy



- Loss



\* Small batches are the reason of these noisy fluctuations "high variance"

\* Can not do large batches, for memory reasons

\* do some smoothing for e.g. exp avg

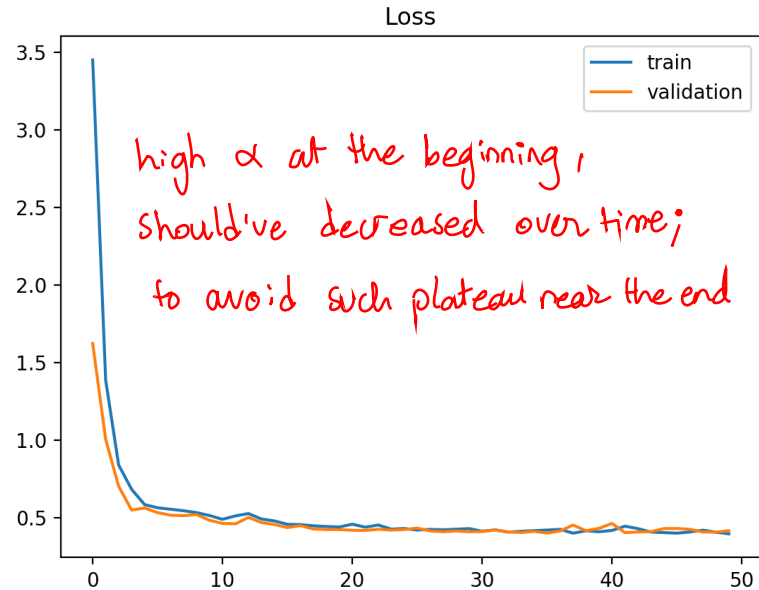

NB

One good practice is to decrease  $\alpha$ , once loss starts to decrease slowly.

Sometimes, however, iterating over more epochs would be ideal.

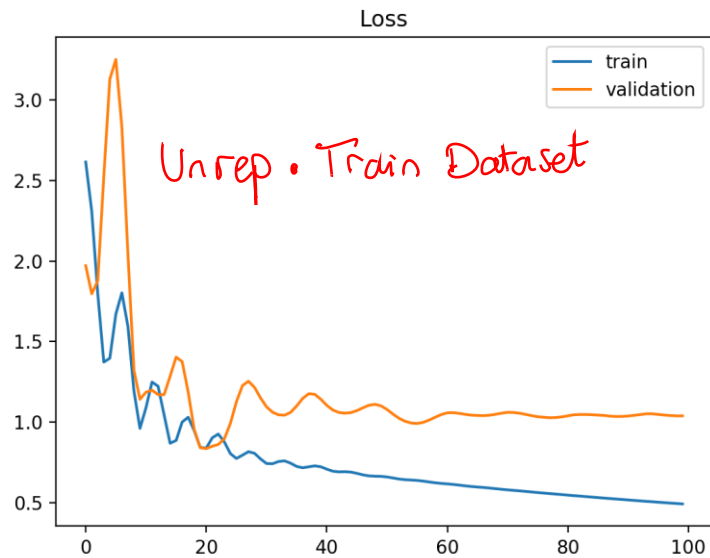
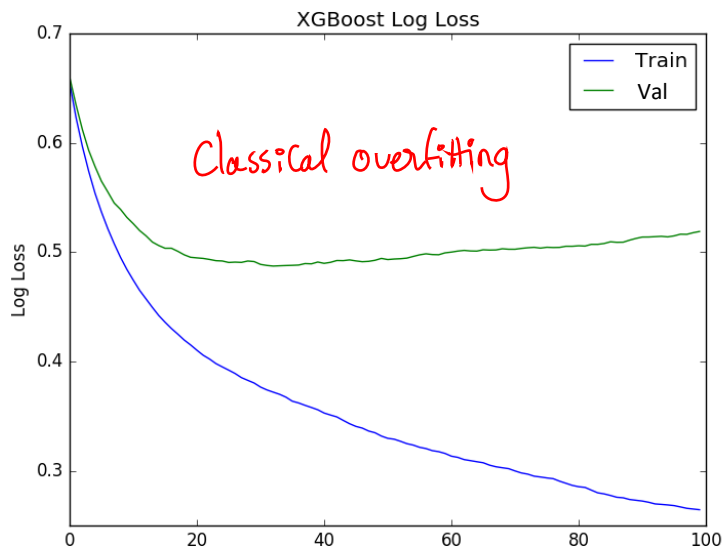


Typically,  $\alpha$  shall be high at the beginning, then decrease over time.


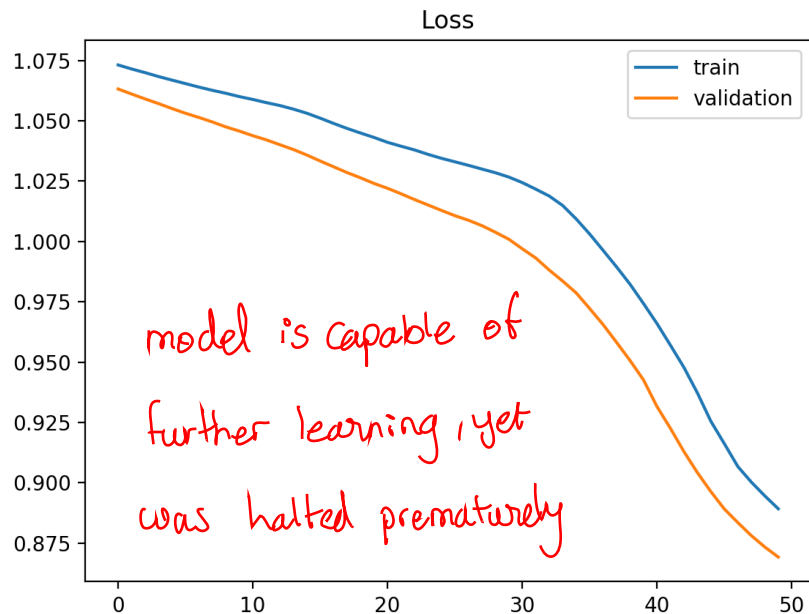


Source: <https://machinelearningmastery.com/learning-curves-for-diagnosing-machine-learning-model-performance/>

# Overfitting Curves



# Other Curves



Underfitting (loss still decreasing)

- ↳ flat loss regardless of training
- ↳ training continues to decrease until the end

Validation Set is easier than Training set

# To Summarize

Important

- Underfitting
  - Training and validation losses decrease even at the end of training
- Overfitting
  - Training loss decreases and validation loss increases
- Ideal Training
  - Small gap between training and validation loss, and both go down at same rate (stable without fluctuations).

## Bad Signs

- Training error not going down
- Validation error not going down
- Performance on validation better than on training set
- Tests on train set different than during training

### • Bad Practice

- Training set contains **test data**
- Debug algorithm on **test data**

Never touch during  
development or  
training

# Hyperparameters



anything but  
model parameters

- Network architecture (e.g., num layers, #weights)
- Number of iterations
- Learning rate(s) (i.e., solver parameters, decay, etc.)
- Regularization (more later next lecture)
- Batch size
- ...

**\*** Overall:  
learning setup + optimization = hyperparameters

# Hyperparameter Tuning

- Methods:
  - **Manual** search:
    - most common 😊
  - **Grid** search (structured, for 'real' applications)
    - Define ranges for all parameters spaces and select points
    - Usually pseudo-uniformly distributed
    - Iterate over all possible configurations
  - **Random** search:
    - Like grid search but one picks points at random in the predefined ranges



ALWAYS

# How to Start

- Start with single training sample
  - Check if output correct *i.e debug loss fn , not the input*
  - Overfit → train accuracy should be 100% because input just memorized
- Increase to handful of samples (e.g., 4)
  - Check if input is handled correctly
- Move from overfitting to more samples
  - 5, 10, 100, 1000, ...
  - At some point, you should see generalization




*no need to extract features,  
just learn the input*





# Find a Good Learning Rate


- Use all training data with small weight decay
- Perform initial loss sanity check e.g.,  $\log(\mathcal{C})$  for softmax with  $\mathcal{C}$  classes
- Find a learning rate that makes the loss drop significantly (exponentially) within 100 iterations
- Good learning rates to try:  $1e-1$ ,  $1e-2$ ,  $1e-3$ ,  $1e-4$



# Coarse Grid Search

usually out of  
compute budget

- Choose a few values of learning rate and weight decay around what worked from
- Train a few models for a few epochs.
- Good weight decay to try:  $1e-4$ ,  $1e-5$ , 0




# Refine Grid

- Pick best models found with coarse grid.
- Refine grid search around these models.
- Train them for longer (10-20 epochs) without learning rate decay
- ✱ Study loss curves <- most important debugging tool!

# Timings

- How long does each iteration take?
  - Get precise timings!
  - If an iteration exceeds **500ms**, things get dicey
- Look for bottlenecks
  - Dataloading: smaller resolution, compression, train from SSD
  - Backprop
- Estimate total time
  - How long until you see some pattern?
  - How long till convergence?



# Network Architecture

- Frequent mistake: *"Let's use this super big network, train for two weeks and we see where we stand."*

- ✱ Instead: start with simplest network possible
  - Rule of thumb divide #layers you started with by 5
- Get debug cycles down
  - Ideally, minutes



# Debugging

- Use train/validation/test curves
  - Evaluation needs to be consistent
  - Numbers need to be comparable

- \* Only make one change at a time
  - “I’ve added 5 more layers and double the training size, and now I also trained 5 days longer. Now it’s better, but why?”

# Common Mistakes in Practice

- Did not overfit to single batch first
- Forgot to toggle train/eval mode for network
  - Check later when we talk about dropout...
- Forgot to call .zero\_grad() (*in PyTorch*) before calling **.backward()**
- ✗ Passed softmaxed outputs to a loss function that expects raw logits


# Tensorboard:

## Visualization in


## Practice




# TensorBoard: Compare Train/Val Curves




# TensorBoard: Compare Different Runs



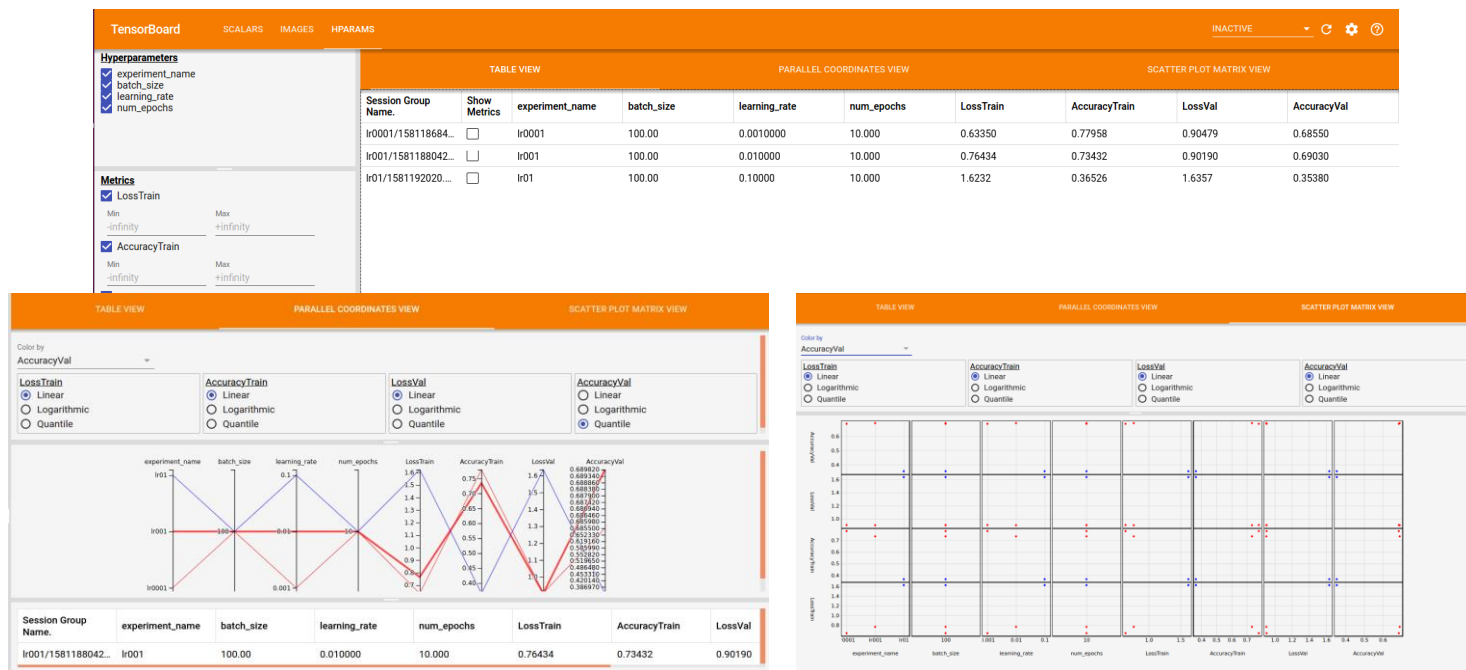
# TensorBoard: Visualize Model Predictions



# TensorBoard: Visualize Model Predictions



# TensorBoard: Compare Hyperparameters



# Next Lecture

- Next lecture
  - More about training neural networks: output functions, loss functions, activation functions
- Check the exercises 😊

See you next week 😊

# References

- Goodfellow et al. "Deep Learning" (2016),
  - Chapter 6: Deep Feedforward Networks
- Bishop "Pattern Recognition and Machine Learning" (2006),
  - Chapter 5.5: Regularization in Network Nets
- <http://cs231n.github.io/neural-networks-1/>
- <http://cs231n.github.io/neural-networks-2/>
- <http://cs231n.github.io/neural-networks-3/>