


Introduction to Deep Learning (I2DL)

Exercise 8: Autoencoder

Today's Outline

- Exam
 - Mock Exam
- Hyperparameter tuning
- Exercise 8
 - Batch Normalization & Dropout
 - Transfer Learning
 - Autoencoder
- Personal: Github/Exposure



Mock Exam & Update

- Exam Structure via Mock Exam

- Multiple Choice
- Written Questions
- No coding

- Exam

- No concrete information from the university
- We will keep you posted
- Additional questions after CNNs and once we have more info

Introduction to Deep Learning (I2DL)

Mock Exam

IN2346 - SoSe 2020

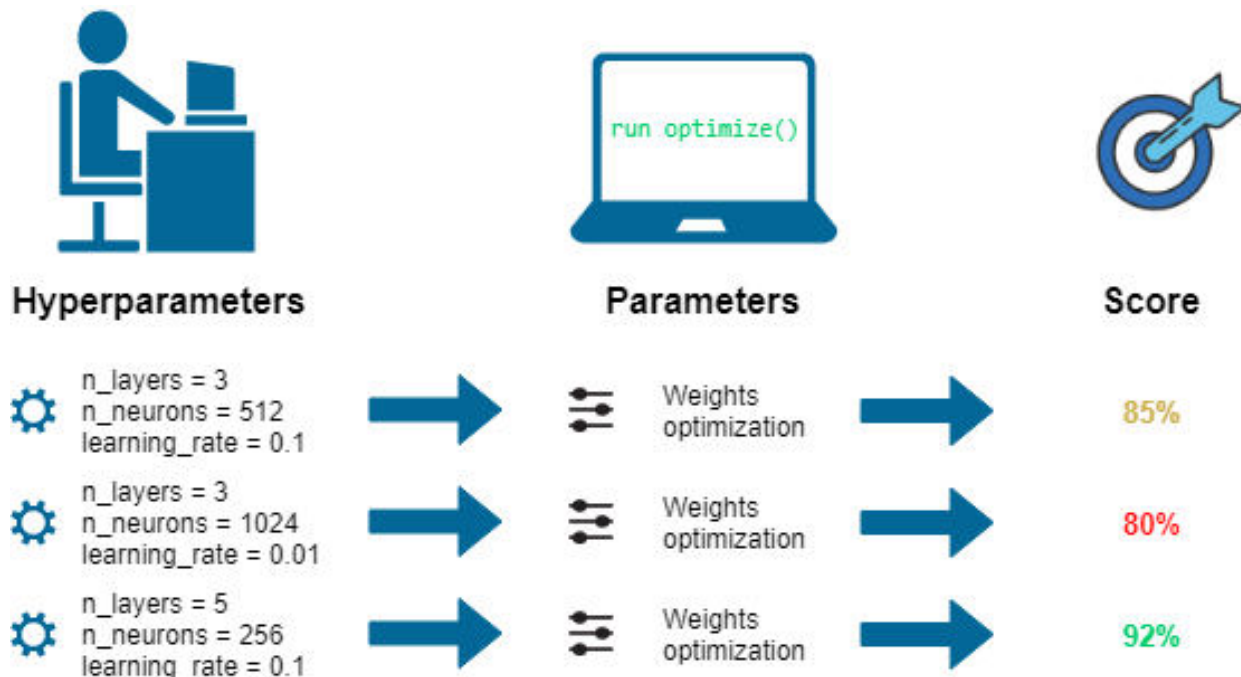
Technical University of Munich

	Problem	Full Points	Your Score
1	Multiple Choice	10	
2	Short Questions	12	
3	Backpropagation	9	
Total		31	

Total Time: **31 Minutes**

Allowed Ressources: **None**

Hyperparameter Tuning



Source: <https://images.deepai.org/glossary-terms/05c646fe1676490aa0b8cab0732a02b2/hyperparams.png>

Hyperparameter Tuning

- Slides on Piazza
 - Check them out if you haven't done it yet




- It is important
 - Regardless of your resources
 - There is no all in one recipe
 - If you need more practice: optional submission on CIFAR10 (Solutions will be discussed next week)

Improve your
training!


Batch Normalization

- All we want is that our activations do not die out



Batch Normalization

- Wish: Unit Gaussian activations



Mean of your mini-batch examples over feature k

Unit gaussian

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{\text{Var}[\mathbf{x}^{(k)}]}}$$

Batch Normalization

- 1. Normalize

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

- 2. Allow the network to change the range

$$\mathbf{y}^{(k)} = \gamma^{(k)} \hat{\mathbf{x}}^{(k)} + \beta^{(k)}$$


backprop

The network *can* learn to undo the normalization


$$\gamma^{(k)} = \sqrt{Var[\mathbf{x}^{(k)}]}$$

$$\beta^{(k)} = E[\mathbf{x}^{(k)}]$$

Dropout



(a) Standard Neural Net






(b) After applying dropout.

Forward ↑

- Using half the network = half capacity

Transfer Learning

Transfer Learning: Example Scenario



- Need to build a Cat classifier
- Only have a few images ~10 000

Transfer Learning

- Problem Statement:
 - Training a Deep Neural Network needs a lot of data
 - Collecting much data is expensive or just not possible
- Idea:
 - Some problems/ tasks are closely related
 - Can we transfer knowledge from one task to another?
 - Can we re-use (at least parts of) a pre-trained network for the new task?


Transfer Learning

Distribution



Large dataset

Distribution



Small dataset



Use what has been
learned for another
setting




Transfer Learning




Coloring Legend:

 Untrained

 Trained



Transfer Learning




Coloring Legend:

 Untrained

 Trained

Transfer Learning



Coloring Legend:

 Untrained

 Trained

Feature Extractor

Classifier


Maybe freeze weights/
slower learning rate/
nothing special

Newly initialized
head

Application: Autoencoder (Sub 8)


Autoencoder

- Task
 - Reconstruct the input given a lower dimensional bottleneck
 - Loss: L1/L2 per pixel
- Actually need no labels!
- Without non-linearities: similar to PCA



Transfer Using an Autoencoder

- Step 1:
 - Train an Autoencoder on a large (maybe unlabeled) dataset very similar to your target dataset
- Step 2:
 - Take pre-trained Autoencoder and use it as the first part of a classification architecture for your target dataset



Personal Note: Github/Exposure

- Posting I2DL solutions is not a helpful git for you
 - Maybe among other students...
- What is useful?
 - Something to talk about in interviews
- Projects:
 - internships/ guided research/ any task basically
 - Document your process (blog), show and visualize your data processing, discuss design decisions and publish code

See you next week!