

Grundlagen der künstlichen Intelligenz – Constraint Satisfaction Problems

Matthias Althoff

TU München

October 31, 2019

Organization

- 1 Defining Constraint Satisfaction Problems
- 2 Backtracking Search for Constraint Satisfaction Problems
- 3 Heuristics for Backtracking Search
-  4 Interleaving Search and Inference
- 5 The Structure of Constraint Satisfaction Problems

The content is covered in the AI book by the section “Constraint Satisfaction Problems”.

Learning Outcomes

- You can explain the difference between constraint satisfaction problems (CSPs) and standard search problems.
- You can judge whether a problem is a CSP.
- You can create formally defined CSPs from a problem description.
- You can create constraint graphs.
- You can apply backtracking search.
- ✗ You can apply and decide when to use the following heuristics: *Minimum Remaining Values, Degree Heuristic, and Least Constraining Value*.
- ✗ You can apply techniques interleaving search and inference: *Forward Checking and Arc Consistency Algorithm*.
- You can exploit the structure of CSPs and reduce the complexity of solving tree-structured and nearly tree-structured CSPs.

Difference to Standard Search Problems

Standard Search Problems

Each **state** is atomic, or indivisible, and has no internal structure.

Constraint Satisfaction Problems (CSPs)

- We use a factored representation of each state: a set of variables, each of which has a value.
- The goal test is whether each variable has a value that satisfies all constraints of the problem.

Benefit: Allows useful general-purpose algorithms with more power than standard search algorithms by exploiting the structure of the states.

Set-Up

Constraint Satisfaction Problem

A constraint satisfaction problem is a tuple (X, D, C) , where:

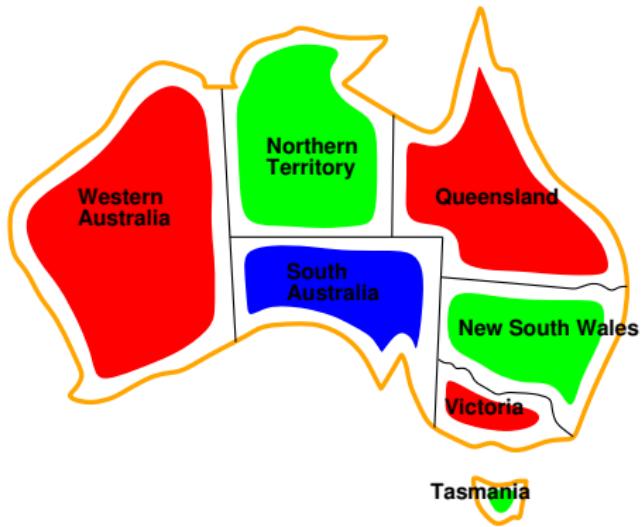
- $X = \{X_1, \dots, X_n\}$ is a set of variables,
 - $D = \{D_1, \dots, D_n\}$ is a set of the respective domains of values, and
 - $C = \{C_1, \dots, C_m\}$ is a set of constraints.
-
- Each domain D_i consists of a set of allowable values $\{v_1, \dots, v_k\}$ for variable X_i .
 - Each constraint C_i consists of a pair $\langle \text{scope}, \text{rel} \rangle$, where scope is a tuple of variables that participate in the constraint and rel is a relation that defines the possible values.

Example Problem: Map Coloring 'Four Color Theorem'



- **Variables:** $X = \{WA, NT, Q, NSW, V, SA, T\}$
- **Domains:** $D_i = \{\text{red}, \text{green}, \text{blue}\}$
- **Constraints:** adjacent regions must have different colors
 $C = \{SA \neq WA, SA \neq NT, SA \neq Q, SA \neq NSW, SA \neq V, WA \neq NT,$
 $NT \neq Q, Q \neq NSW, NSW \neq V\}; \quad (SA \neq WA \text{ is short for } ((SA, WA), SA \neq WA))$

Possible Solution of Map Coloring



Solutions are assignments satisfying all constraints, e.g.,

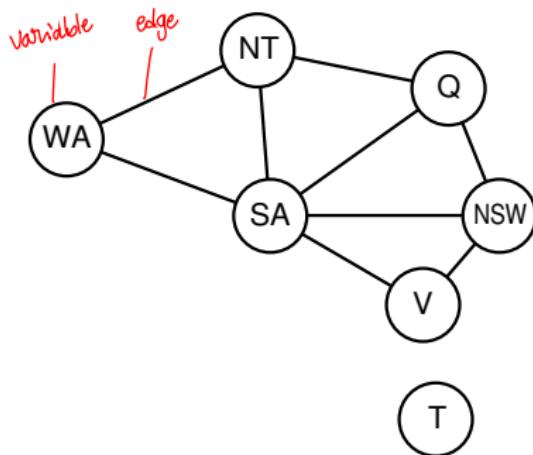
$\{WA = \text{red}, NT = \text{green}, Q = \text{red}, NSW = \text{green}, V = \text{red}, SA = \text{blue}, T = \text{green}\}$

Constraint Graph

It can be helpful to visualize a constraint satisfaction problem as a **constraint graph**:

- Nodes correspond to variables,
- edges connect two variables that participate in a constraint.

variables
 values
 3^6 assignments
 Since it has the most neighbors



Standard search would require searching all combinations, but

- by fixing e.g., $SA = \text{blue}$, none of the five neighbors can choose blue . This reduces from $3^5 = 243$ assignments in standard search to only $2^5 = 32$ (reduction by 87%).
- Tasmania is even an independent subproblem.

Varieties of Constraint Satisfaction Problems

Discrete domains

Finite domains; size $d \Rightarrow \mathcal{O}(d^n)$ complete assignments

- For instance, Boolean CSPs, incl. Boolean satisfiability (NP-complete).

Infinite domains (integers, strings, etc.)

- For instance, job scheduling, variables are start/end days for each job;
- requires a constraint language, e.g., $StartJob_1 + 5 \leq StartJob_3$;
- **linear** constraints solvable, **nonlinear** undecidable.

Continuous domains (not part of this lecture)

- For instance, start/end times for Hubble Telescope observations.
- Linear constraints solvable in polynomial time by linear programming methods.

Varieties of Constraints

- **Unary** constraints involve a single variable,
e.g., $SA \neq green$
- **Binary** constraints involve pairs of variables,
e.g., $SA \neq WA$
- **Higher-order** constraints involve 3 or more variables,
e.g., $SA \neq WA \neq NT$. Higher order constraints can be rewritten as several binary constraints. Previous example: $SA \neq WA$ and $WA \neq NT$.

For that reason, we only consider binary constraints from now on.

- **Preferences** (soft constraints), e.g., *red* is better than *green* is often representable by a cost for each variable assignment
→ constrained optimization problems (not part of this lecture)

Real-World Constraint Satisfaction Problems

- **Assignment problems**, e.g., who teaches what class?
- **Timetabling problems**, e.g., which class is offered when and where?
- **Hardware configuration**, e.g., what kind of processor, memory, bus system, motherboards, etc., can be combined?
- **Spreadsheets**, e.g., check constraints to ensure correctness of data.
- **Transportation scheduling**, e.g., scheduling of trains so that changing trains is easy.
- **Factory scheduling**, e.g., determining in which order pieces have to be assembled.
- **Floorplanning**, e.g., how should the rooms in a house be arranged, when the living room should face south and the bathroom should have a window to the outside?

Many real-world problems involve real-valued variables.

Standard Search Formulation

Let's start with the naive search approach, then fix it.

States are defined by the values assigned so far:

- **Initial state:** the empty assignment. \emptyset
- **Successor function:** assign a value to an unassigned variable not conflicting with the current assignment.
- **Goal test:** the current assignment is complete.

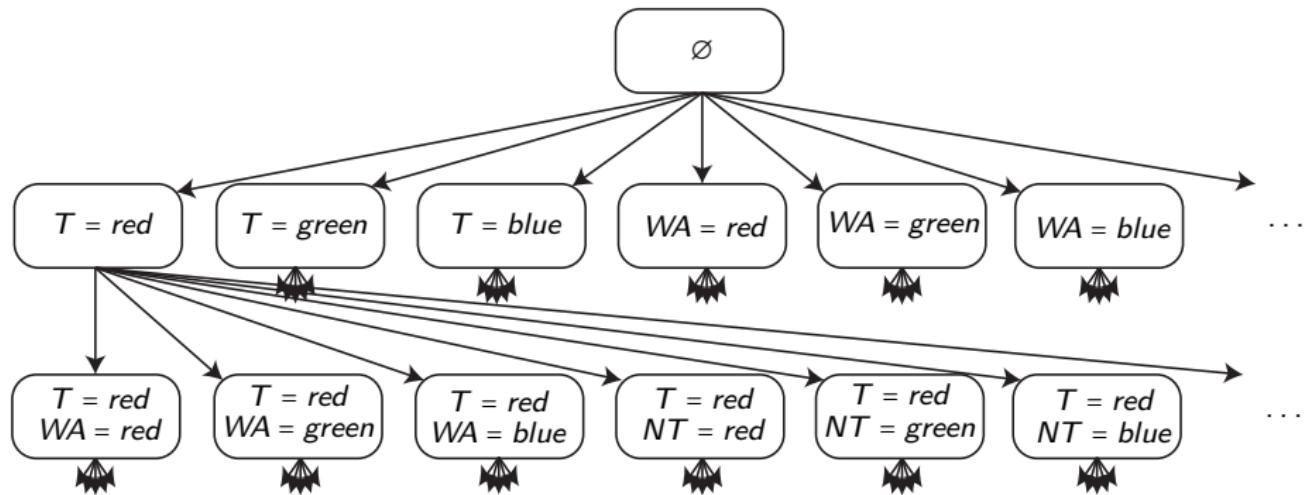
Comments:

- ① This is the same for all CSPs! ☺
- ② Every solution appears at depth n with n variables
⇒ use depth-first search.
- ③ Path is irrelevant.
- ④ $b = (n - l)d$ at depth l , hence $n!d^n$ leaves in the worst case!!!! ☺
(n : nr. of variables, d : nr. of values; assumption: all variables have same nr. of values)

very

Example of the Naive Approach

We use the map coloring problem:



Number of nodes in the worst case:

- **First level:** $n \cdot d$ nodes (n : number of variables, d : number of values).
- **Second level:** $(n \cdot d)((n - 1) \cdot d)$ nodes.
- **n^{th} level:** $\prod_{l=0}^{n-1} (n - l) \cdot d = \prod_{l=0}^{n-1} (n - l) \cdot \prod_{l=0}^{n-1} d = n!d^n$ nodes.

Backtracking Search

- We can drastically improve the naive approach by considering that variable assignments are commutative:

$[WA = \text{red} \text{ then } NT = \text{green}]$ same as $[NT = \text{green} \text{ then } WA = \text{red}]$

- ⇒ Only consider assignments to a single variable at each node
- ⇒ $b = d$ and there are only d^n leaves

* Depth-first search for CSPs with single-variable assignments is called **backtracking** search.

- Backtracking search is the basic uninformed algorithm for CSPs (e.g., solves the n -queens problem for $n \approx 25$).

NB : it can be applied only for problems which admits the concept of a "partial candidate soln" and a relatively quick test of whether it can possibly be completed to a valid soln.

Backtracking Search: Map-Coloring Example (1)

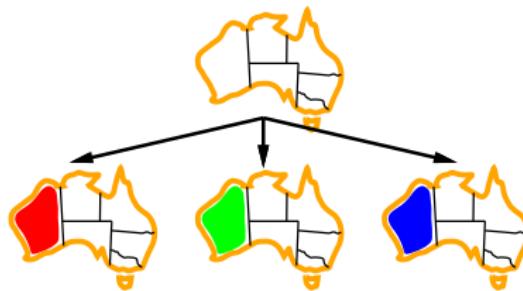
Part of the search tree for the map-coloring problem:



* Backtracking Search: is used for depth-first search that chooses a value for one variable at a time & backtracks when a variable has no legal values left to assign

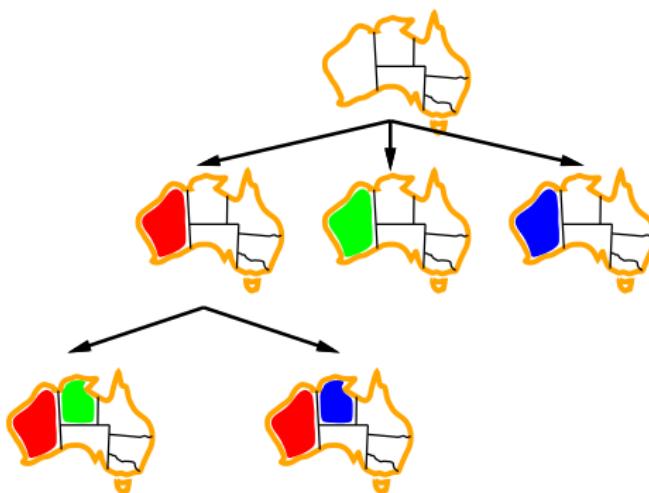
Backtracking Search: Map-Coloring Example (2)

Part of the search tree for the map-coloring problem:



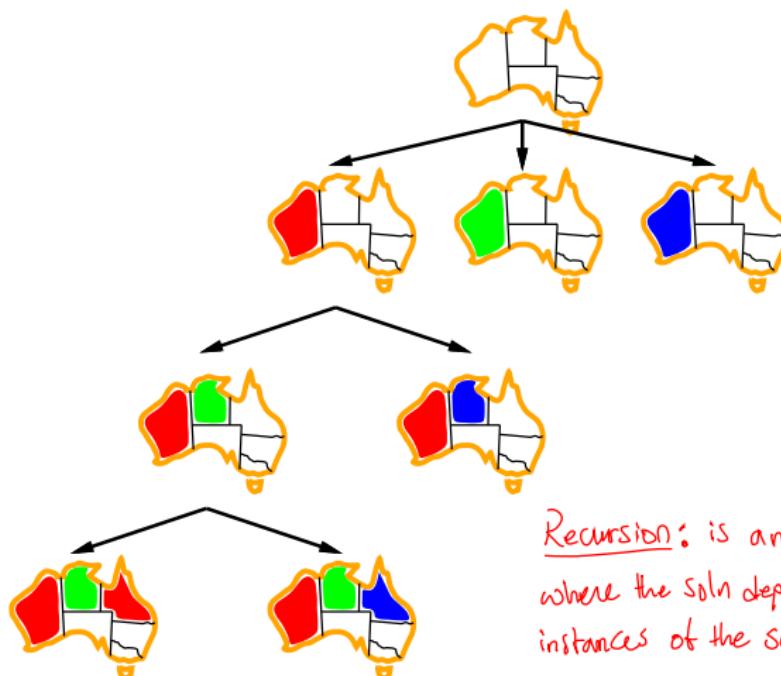
Backtracking Search: Map-Coloring Example (3)

Part of the search tree for the map-coloring problem:



Backtracking Search: Map-Coloring Example (4)

Part of the search tree for the map-coloring problem:



Recursion: is a method of solving a problem where the soln depends on solns to a smaller instances of the same problem.

Backtracking Search: Algorithm

Check: wiki - Backtracking
 'Description of the method'

```
function Backtracking-Search (csp) returns solution/failure
```

```
return Recursive-Backtracking({ }, csp)
```

```
function Recursive-Backtracking (assignment, csp) returns sol./failure
```

```
if assignment is complete then return assignment
```

```
var  $\leftarrow$  Select-Unassigned-Variable(csp)
```

```
for each value in Order-Domain-Values(var, assignment, csp) do
```

```
if value is consistent with assignment given Constraints[csp] then
```

```
    add {var = value} to assignment
```

```
    inferences  $\leftarrow$  Inference(csp, var, value)
```

```
    if inferences  $\neq$  failure then
```

```
        add inferences to assignment
```

```
        result  $\leftarrow$  Recursive-Backtracking(assignment, csp)
```

```
        if result  $\neq$  failure then return result
```

```
    remove {var = value} and inferences from assignment
```

```
return failure
```

Backtracking Search: Heuristics (1)

```
function Recursive-Backtracking (assignment, csp) returns sol./failure
  if assignment is complete then return assignment
  var ← Select-Unassigned-Variable(csp)
  for each value in Order-Domain-Values(var, assignment, csp) do
    if value is consistent with assignment given Constraints[csp] then
      add {var = value} to assignment
      inferences ← Inference(csp,var,value)
      if inferences ≠ failure then
        add inferences to assignment
        result ← Recursive-Backtracking(assignment,csp)
        if result ≠ failure then return result
      remove {var = value} and inferences from assignment
  return failure
```

Important heuristics (domain-specific heuristics as for A* search often not required):

- ① Which variable should be assigned next?

Backtracking Search: Heuristics (2)

```

function Recursive-Backtracking (assignment, csp) returns sol./failure
  if assignment is complete then return assignment
  var  $\leftarrow$  Select-Unassigned-Variable(csp)
  for each value in Order-Domain-Values(var, assignment, csp) do
    if value is consistent with assignment given Constraints[csp] then
      add  $\{var = value\}$  to assignment
      inferences  $\leftarrow$  Inference(csp, var, value)
      if inferences  $\neq$  failure then
        add inferences to assignment
        result  $\leftarrow$  Recursive-Backtracking(assignment, csp)
        if result  $\neq$  failure then return result
      remove  $\{var = value\}$  and inferences from assignment
  return failure

```

Important heuristics (domain-specific heuristics as for A* search often not required):

- ① Which variable should be assigned next?
- ② In what order should its values be tried?

Backtracking Search: Heuristics (3)

```

function Recursive-Backtracking (assignment, csp) returns sol./failure
  if assignment is complete then return assignment
  var  $\leftarrow$  Select-Unassigned-Variable(csp)
  for each value in Order-Domain-Values(var, assignment, csp) do
    if value is consistent with assignment given Constraints[csp] then
      add  $\{var = value\}$  to assignment
      inferences  $\leftarrow$  Inference(csp, var, value)
      if inferences  $\neq$  failure then
        add inferences to assignment
        result  $\leftarrow$  Recursive-Backtracking(assignment, csp)
        if result  $\neq$  failure then return result
      remove  $\{var = value\}$  and inferences from assignment
  return failure

```

Important heuristics (domain-specific heuristics as for A* search often not required):

- ① Which variable should be assigned next?
- ② In what order should its values be tried?
- ③ Can we detect inevitable failure early?

Backtracking Search: Heuristics (4)

function Recursive-Backtracking (*assignment*, *csp*) **returns** sol./failure

if *assignment* is complete **then return** *assignment*

var \leftarrow Select-Unassigned-Variable(*csp*)

for each *value* in Order-Domain-Values(*var*, *assignment*, *csp*) **do**

if *value* is consistent with *assignment* given Constraints[*csp*] **then**

 add {*var* = *value*} to *assignment*

inferences \leftarrow Inference(*csp*, *var*, *value*)

if *inferences* \neq failure **then**

 add *inferences* to *assignment*

result \leftarrow Recursive-Backtracking(*assignment*, *csp*)

if *result* \neq failure **then return** *result*

 remove {*var* = *value*} and *inferences* from *assignment*

return failure



Important heuristics (domain-specific heuristics as for A* search often not required):

① Which variable should be assigned next? SA: most connected

② In what order should its values be tried? red, green, ...

③ Can we detect inevitable failure early? Inference methods

④ Can we take advantage of problem structure? (see slide 44 onwards)

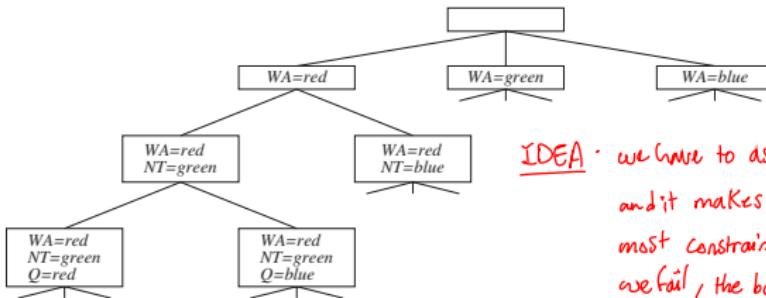
Minimum Remaining Values

The backtracking algorithm contains $\text{var} \leftarrow \text{Select-Unassigned-Variable}(csp)$.

- The simplest strategy is to choose the next unassigned variable in order ($\{X_1, X_2, \dots\}$).
- The above strategy seldomly is the most efficient one.

Example (see figure): for $WA = \text{red}$ and $NT = \text{green}$, there is only one possibility for SA , so it makes sense to assign $SA = \text{blue}$ rather than assigning Q . After SA is assigned, the choices for Q , NSW , and V are all forced.

The intuitive idea of choosing the variable with the fewest possible values first is called **minimum-remaining-values** (MRV) heuristic.

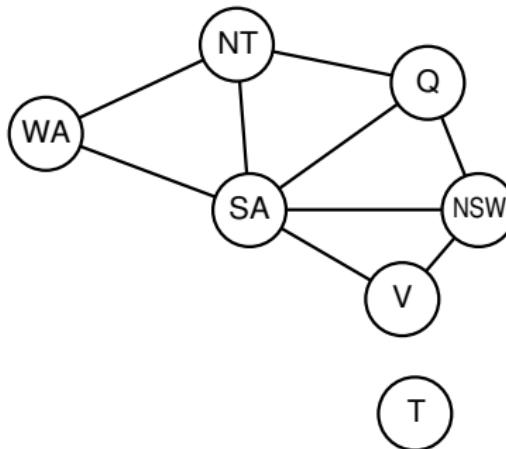


IDEA : we have to assign all variables eventually and it makes more sense to pick the most constrained values now so that if we fail, the backtracking is not expensive

Degree Heuristic

- The MRV heuristic does not help in choosing the first region in Australia.
- A good choice is to select the variable that is involved in the largest number of constraints on other unassigned variables, called **degree heuristic**.

Example (see figure): SA is involved in 5 constraints, while other variables are only in 2 or 3 constraints, except for T. Once SA is chosen, continuing use of degree heuristic solves the problem without any false step.



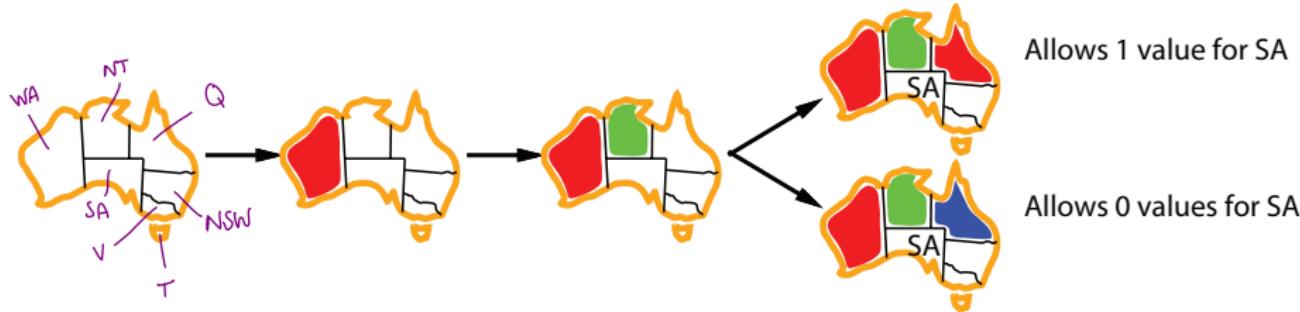
Least Constraining Value

The backtracking algorithm contains

for each value in Order-Domain-Values(*var, assignment, csp*) **do**.

- Optimal order of choosing values?
- A good choice is to select the value that rules out the fewest choices for neighboring values in the constraint graph, called **least-constraining-value** heuristic.

Example (see figure): WA = red and NT = green are selected and our next choice is Q. Blue would be a bad choice since it removes all options for the neighbor SA, while red leaves an option for SA. Thus, red is preferred.



Comments on Variable and Value Selection



Why should variable selection be fail-first, but value selection be fail-last?

- For a wide variety of problems, choosing variables with the minimum number of remaining values helps to prune the search tree (in the end, each variable has to be selected anyways).
- For value ordering, the trick is that we only need one solution; therefore, it makes sense to look at the most likely values first.

values that least strict the neighbors

Overall, we present heuristics, so it is not guaranteed that the proposed heuristics always provide fast solutions. But they work well in most cases!

Inference in Constraint Satisfaction Problems

Inference

Act or process of deriving logical conclusions from known premises.

Inference can be applied at different times during the backtracking algorithm:

- **After each assignment:** see Inference() in backtracking algorithm.
- **As pre-processing:** before applying the backtracking algorithm.

We look at the two following inference techniques (others exist):

- 1 • **Forward checking** (after each assignment): inconsistent values of only neighboring variables are removed.
- 2 • **Arc consistency algorithm** (after each assignment or as pre-processing): inconsistent values of all variables are removed.

Definition of Arc Consistency

Cornell_CSP, slide 20

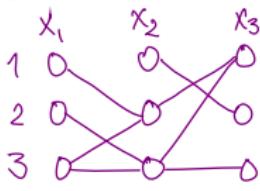
* Constraint propagation (e.g arc-consistency) does additional work to constraint values & detect inconsistencies

- **Arc consistency of a variable:** variable X_i is arc-consistent with variable X_j , if for every value in the domain D_i there exists a value in D_j satisfying the binary constraint of the arc (X_i, X_j) .

Example: X is arc-consistent with Y for the constraint $Y = X^2$ if $D_X = \{0, 1, 2, 3\}$ and $D_Y = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$, but Y is not arc-consistent with X (direction of the arc matters).

- * **Arc consistency of a CSP:** a constraint graph is arc-consistent if every variable is arc-consistent with every other variable.

Example: the constraint graph $Y = X^2$ is arc-consistent, if $D_X = \{0, 1, 2, 3\}$ and $D_Y = \{0, 1, 4, 9\}$.



X_2 is arc-consistent with X_3 but not with X_1 , as the value $X_2=1$ does not correspond to any value for X_1 .

Forward Checking (1) → prevents assignments that guarantee later failure

After a variable X_i is assigned during backtracking search, forward checking makes all variables X_j constrained with X_i arc-consistent with X_i .

Initial domains:

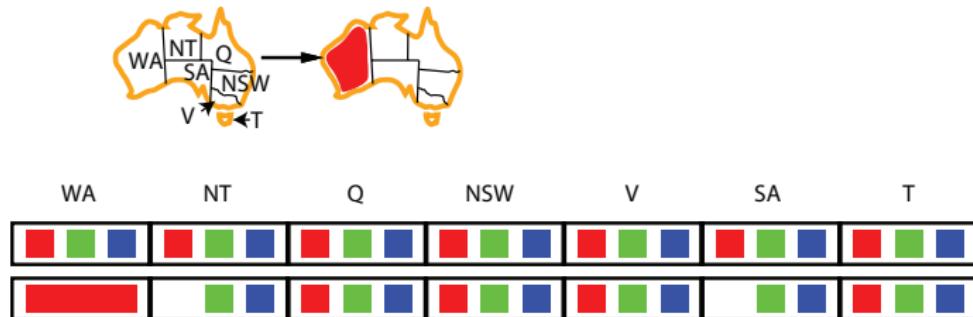


WA	NT	Q	NSW	V	SA	T
Red	Green	Blue	Red	Green	Blue	Red

Forward Checking (2)

After a variable X_i is assigned during backtracking search, forward checking makes all variables X_j constrained with X_i arc-consistent with X_i .

After assigning $WA = red$:



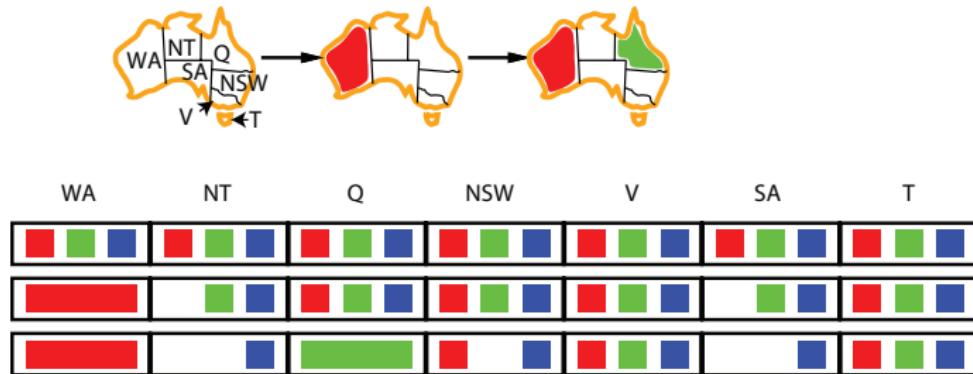
* **Comment:** forward checking removes the value *red* from all neighbors of *WA* so that they are arc-consistent with *WA*. Thus, *NT* is arc-consistent with *WA*, and *SA* is arc-consistent with *WA*.

e.g. X arc cons with Y
 $\{no red\} \rightarrow \{red\}$

Forward Checking (3)

After a variable X_i is assigned during backtracking search, forward checking makes all variables X_j constrained with X_i arc-consistent with X_i .

After assigning $Q = \text{green}$:

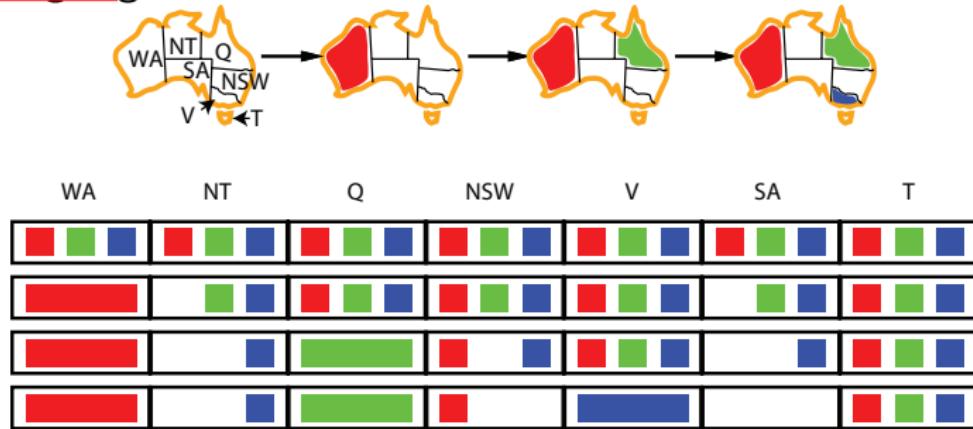


* **Comment:** forward checking removes the value *green* from all neighbors of Q .
Note that forward checking does not change the domains of other variables and does not assign variables whose domain only contains a single value.

Forward Checking (4)

After a variable X_i is assigned during backtracking search, forward checking makes all variables X_j constrained with X_i arc-consistent with X_i .

After assigning $V = \text{blue}$:



Comment: forward checking results in a failure, since the domain for SA becomes empty. Thus, the backtracking search algorithm will backtrack: it revokes the last inconsistent assignment $V = \text{blue}$.

Arc Consistency Algorithm

function AC-3 (*csp*, *queue*) **returns** failure or the reduced *csp* otherwise

inputs: *csp*: a binary CSP, *queue*: a queue of arcs (X_i, X_j)

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

⇒ AC-3 is the most popular algorithm for arc-consistency

if Remove-Inconsistent-Values(X_i, X_j) **then**

if size of Domain(X_i) = 0 **then return** failure

for each X_k **in** Neighbors[X_i] \{ X_j \} **do**

 add (X_k, X_i) to *queue*

return *csp*

function Remove-Inconsistent-Values (X_i, X_j) **returns** true iff succeeds

removed \leftarrow false

for each x **in** Domain[X_i] **do**

if no value y in Domain[X_j] allows (x, y) to satisfy the constraint of (X_i, X_j)
then delete x from Domain[X_i]; *removed* \leftarrow true

return *removed*

* Time complexity is $\mathcal{O}(cd^3)$ (*c*: number of arcs, *d*: maximum domain size).

Initialization of the Arc Consistency Algorithm

The input variable *queue* is initialized depending on when AC-3 is used:

- **after each assignment:** after assigning variable X_i , add the arcs (X_j, X_i) to *queue*, where X_j are all unassigned neighbors of X_i .
- **as pre-processing:** add all arcs of the CSP to *queue*.



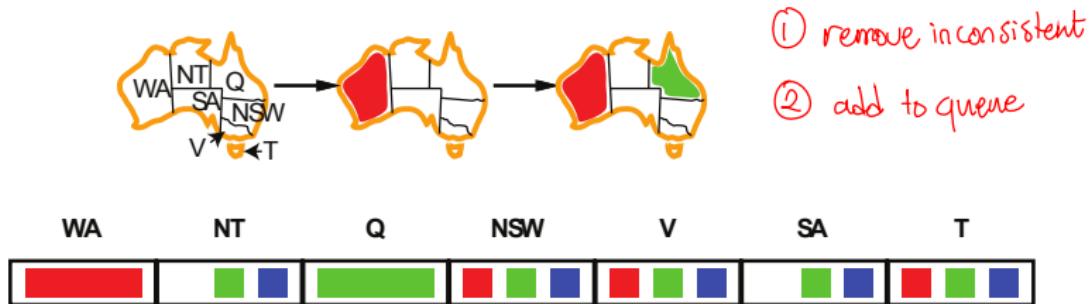
If the arc consistency algorithm terminates successfully, all variables of the CSP are arc-consistent with each other and possibly have reduced domains.

Remember the idea is to form an arc-consistent graph
for CSP, where each variable is arc-consistent
with every other variable

Arc Consistency Algorithm: Example (1)

Example for applying the arc consistency algorithm (AC-3) after an assignment.

In the initial map-coloring problem, we assigned $WA = \text{red}$, applied the arc consistency algorithm in $\text{Inference}(csp, var, value)$, and just assigned $Q = \text{green}$.



* Now, we again apply Inference($csp, var, value$).

- add $(NSW, Q), (SA, Q), (NT, Q)$ to *queue*: since Q has just been assigned, the *queue* is initialized with all arcs to neighbors of Q
- call $\text{AC-3}(csp, queue)$

Arc Consistency Algorithm: Example (2)

function AC-3 (*csp, queue*) **returns** failure or the reduced *csp* otherwise

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

if Remove-Inconsistent-Values(X_i, X_j) **then**

if size of Domain(X_i) = 0 **then return** failure

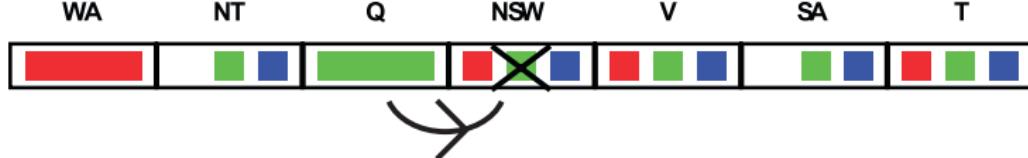
for each X_k **in** Neighbors[X_i] \{ X_j \} **do**

 add (X_k, X_i) to *queue*

Originally, they are
not arc consistent

- $(NSW, Q) \leftarrow \text{Remove-First}(\text{queue})$ *1st neighbor with Q*
- Remove-Inconsistent-Values(*NSW, Q*) **returns** true: green removed from domain of *NSW* *Start*
- add (*V, NSW*), (*SA, NSW*) to *queue*: neighbors of *NSW* (except *Q*)

So that they
are arc
consistent



already
examined

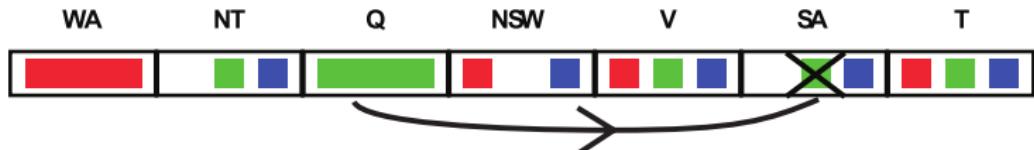
Arc Consistency Algorithm: Example (3)

```

function AC-3 (csp, queue) returns failure or the reduced csp otherwise
while queue is not empty do
     $(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$ 
    if Remove-Inconsistent-Values( $X_i, X_j$ ) then
        if size of Domain( $X_i$ ) = 0 then return failure
        for each  $X_k$  in Neighbors[ $X_i$ ] \{ $X_j$ \} do
            add  $(X_k, X_i)$  to queue

```

- $(SA, Q) \leftarrow \text{Remove-First}(\text{queue})$ *2nd neighbor with C*
- Remove-Inconsistent-Values(SA, Q) **returns** true: green removed from domain of SA
- add $(WA, SA), (NT, SA), (NSW, SA), (V, SA)$ to *queue*: neighbors of SA (except Q)



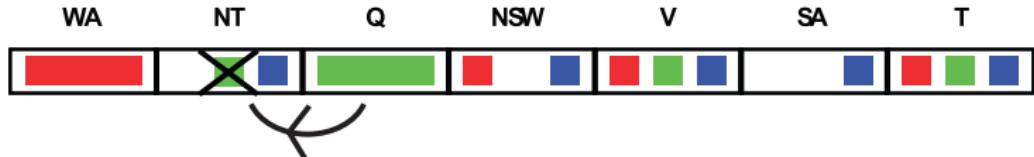
Arc Consistency Algorithm: Example (4)

```

function AC-3 (csp, queue) returns failure or the reduced csp otherwise
while queue is not empty do
     $(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$ 
    if Remove-Inconsistent-Values( $X_i, X_j$ ) then
        if size of Domain( $X_i$ ) = 0 then return failure
        for each  $X_k$  in Neighbors[ $X_i$ ] \{ $X_j$ \} do
            add ( $X_k, X_i$ ) to queue

```

- $(NT, Q) \leftarrow \text{Remove-First}(\text{queue})$ *3rd neighbor to Q*
- Remove-Inconsistent-Values(NT, Q) **returns** true: green removed from domain of NT
- add (WA, NT), (SA, NT) to *queue*: neighbors of NT (except Q)



Arc Consistency Algorithm: Example (5)

function AC-3 (*csp, queue*) **returns** *failure* or the reduced *csp* otherwise

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

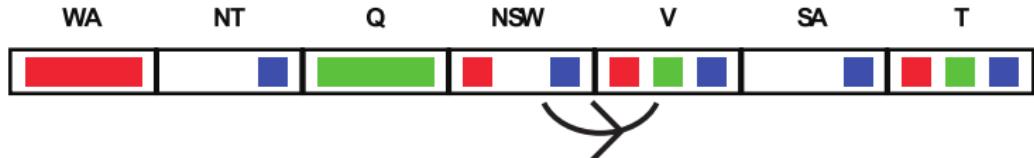
if Remove-Inconsistent-Values(X_i, X_j) **then**

if size of Domain(X_i) = 0 **then return** *failure*

for each X_k **in** Neighbors[X_i] \{ X_j \} **do**

 add (X_k, X_i) to *queue*

- $(V, NSW) \leftarrow \text{Remove-First}(\text{queue})$
 - Remove-Inconsistent-Values($V, \underline{\text{NSW}}$) **returns** false: V is already arc-consistent with *NSW*
- already arc consistent
↗
Start



Arc Consistency Algorithm: Example (6)

function AC-3 (*csp, queue*) **returns** failure or the reduced *csp* otherwise

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

if Remove-Inconsistent-Values(X_i, X_j) **then**

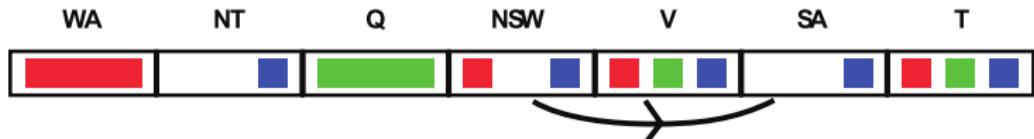
if size of Domain(X_i) = 0 **then return** failure

for each X_k **in** Neighbors[X_i] \{ X_j \} **do**

 add (X_k, X_i) to *queue*

every value is SA
has a value that
satisfies the cons-
traint in NSW

- $(SA, NSW) \leftarrow \text{Remove-First}(\text{queue})$
- Remove-Inconsistent-Values(SA, NSW) **returns** false: SA is already arc-consistent with NSW



Comment: even though SA is arc-consistent with NSW , NSW is not arc-consistent with SA .

Arc Consistency Algorithm: Example (7)

function AC-3 (*csp, queue*) **returns** *failure* or the reduced *csp* otherwise

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

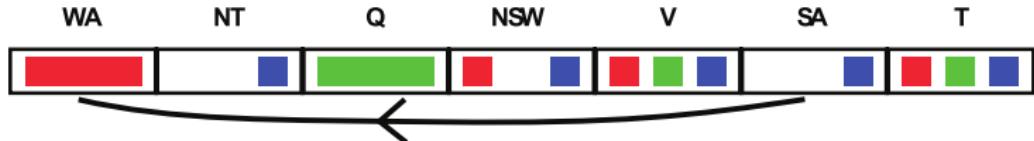
if Remove-Inconsistent-Values(X_i, X_j) **then**

if size of Domain(X_i) = 0 **then return** *failure*

for each X_k **in** Neighbors[X_i] \{ X_j \} **do**

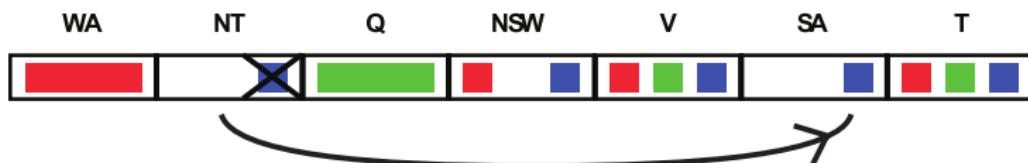
 add (X_k, X_i) to *queue*

- $(WA, SA) \leftarrow \text{Remove-First}(\text{queue})$
- Remove-Inconsistent-Values(*WA, SA*) **returns** *false*: *WA* is already arc-consistent with *SA*



Arc Consistency Algorithm: Example (8)

- $(NT, SA) \leftarrow \text{Remove-First}(queue)$
- Remove-Inconsistent-Values(NT, SA) **returns** true: blue removed from domain of NT
- AC-3 **returns failure**: size of Domain(NT) = 0



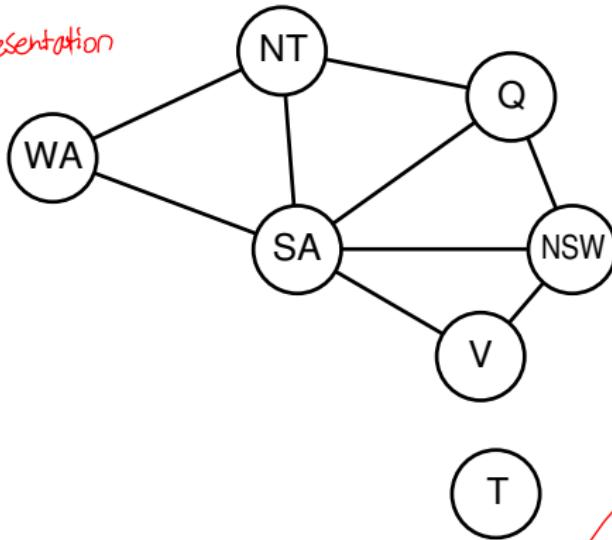
*

Result of applying AC-3: the backtracking search algorithm receives a failure and knows that the CSP cannot be solved with the last assignment $Q = \text{green}$. Thus, the search algorithm backtracks by removing the inconsistent assignment and restoring the modified domains.

Further examples of applying AC-3 are presented in the Exercise.

Problem Structure

* Constraint graph representation
allow analysis of
problem structure



which dramatically
reduces the computational
complexity

Tasmania and the mainland are independent subproblems.

Identifiable as connected components of the constraint graph.

Complexity Reduction of Independent Problems

Completely independent problems provide a fantastic simplification:

- Suppose each subproblem has c variables out of n total.
- Worst-case solution cost is $n/c \cdot d^c$, **linear** in n .
- Worst-case solution cost of the full problem is d^n , **exponential** in n .

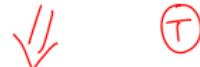
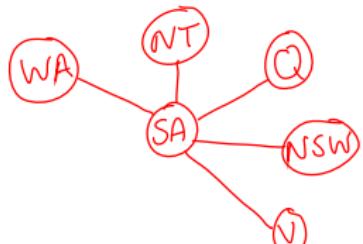
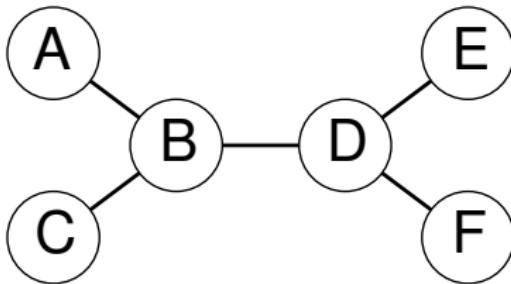
Example

$n = 80$, $d = 2$, $c = 20$:

- full problem: $2^{80} = 4$ billion years at 10 million nodes/sec
(approximate time until earth consumed by the sun)
- independent problems: $4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec

Tree-Structured CSPs

Tree-structured CSPs
can be solved in
linear time

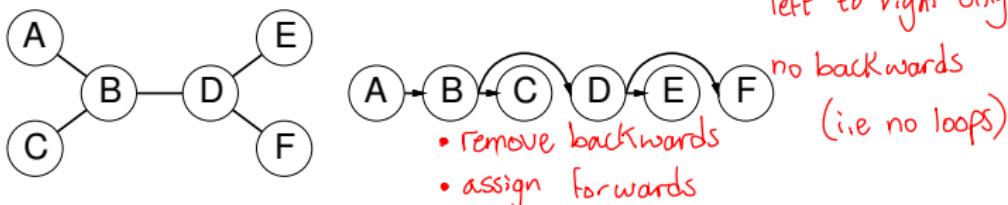


this is no tree structure!
Yet assigning SA to
a value makes it a tree
structure

- If the constraint graph has no loops, the CSP can be solved in $\mathcal{O}(n d^2)$ time (n : nr. of variables, d : domain size).
- Compare to general CSPs, where worst-case time is $\mathcal{O}(d^n)$.
- This property also applies to logical and probabilistic reasoning.

Tree-Structured CSPs: Algorithm

- Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering (**topological sort**).



Any tree with n nodes has only $n - 1$ arcs, so that topological sort is $\mathcal{O}(n)$.

- The obtained variable order X_1, X_2, \dots, X_n is made directly arc-consistent.

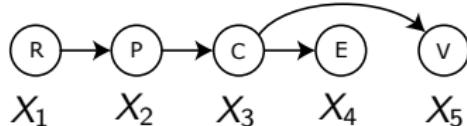
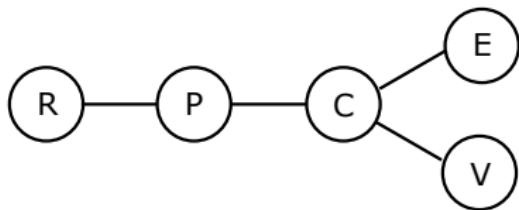
Direct arc consistency

A CSP is direct arc-consistent for the ordered variables X_1, X_2, \dots, X_n if and only if every X_i is arc-consistent with each directly connected X_j for $j > i$.

The complexity is $\mathcal{O}(nd^2)$ since for each node, two variables with d possible domain values have to be compared pairwise.

- Once we have a directed arc-consistent graph, we can march down the list of variables and choose any remaining value; backtracking is not required.

Tree-Structured CSPs: Example (1)



Option 1: The initial domains of all countries are $D_i = \{red, green, blue\}$



$$D_5 = \{red, green, blue\}$$

$$D_4 = \{red, green, blue\}$$

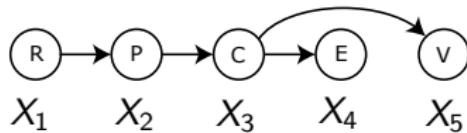
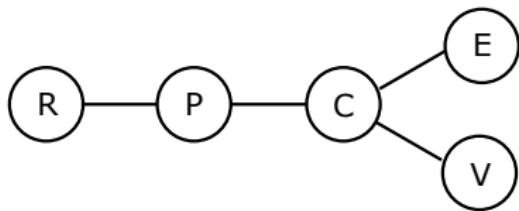
$$D_3 = \{red, green, blue\}$$

$$D_2 = \{red, green, blue\}$$

$$D_1 = \{red, green, blue\}$$

We select $X_1 = red$, $X_2 = green$, $X_3 = blue$, $X_4 = red$, $X_5 = green$.

Tree-Structured CSPs: Example (2)



Option 2: The initial domains of all countries are $D_i = \{ \text{red}, \text{green}, \text{blue} \}$, except that $D_5 = \{ \text{blue} \}$



$$D_5 = \{ \text{blue} \}$$

$$D_4 = \{ \text{red}, \text{green}, \text{blue} \}$$

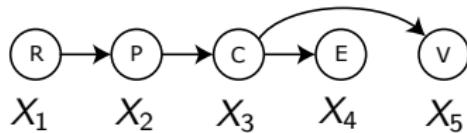
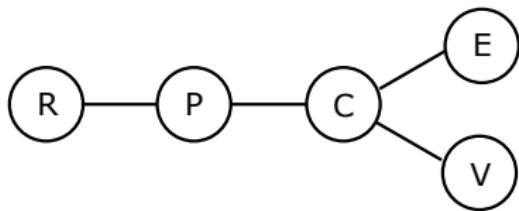
$$D_3 = \{ \text{red}, \text{green} \}$$

$$D_2 = \{ \text{red}, \text{green}, \text{blue} \}$$

$$D_1 = \{ \text{red}, \text{green}, \text{blue} \}$$

We select $X_1 = \text{red}$, $X_2 = \text{green}$, $X_3 = \text{red}$, $X_4 = \text{green}$, $X_5 = \text{blue}$.

Tree-Structured CSPs: Example (3)



Option 3: The initial domains of all countries are $D_i = \{red, green, blue\}$, except that $D_5 = \{blue\}$ and $D_4 = \{green\}$



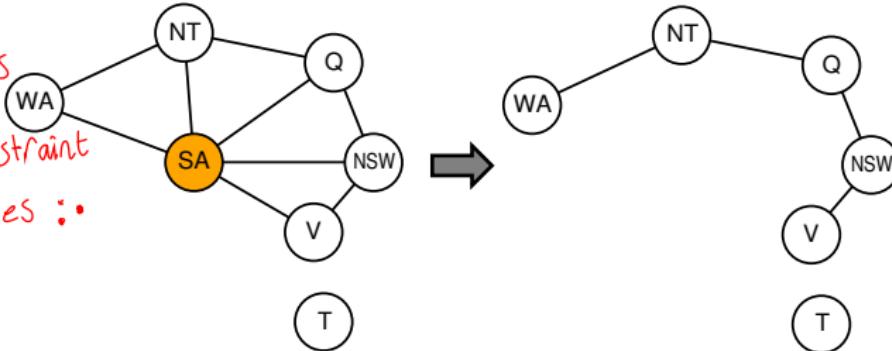
- $D_5 = \{blue\}$
- $D_4 = \{green\}$
- $D_3 = \{red\}$
- $D_2 = \{green, blue\}$
- $D_1 = \{red, green, blue\}$

We select $X_1 = red$, $X_2 = green$, $X_3 = red$, $X_4 = green$, $X_5 = blue$.

Nearly tree-structured CSPs: Conditioning (1)

Conditioning: instantiate a variable, prune its neighbors' domains.

2 approaches
to reduce Constraint
graph to trees ::



In general (as opposed to map coloring), the value chosen for SA could be the wrong one, requiring us to try each possible value:

- ① Choose a subset $S \subset X$ of the CSP's variables X such that the constraint graph becomes a tree after removal of S . [in this case SA]
- ② For each possible constraint-satisfying assignment to variables in S ,
 - ① remove values from the other domains that are inconsistent with S , and
 - ② if the remaining CSP has a solution, return it with the one of S .

Nearly tree-structured CSPs: Conditioning (2)

A few comments:

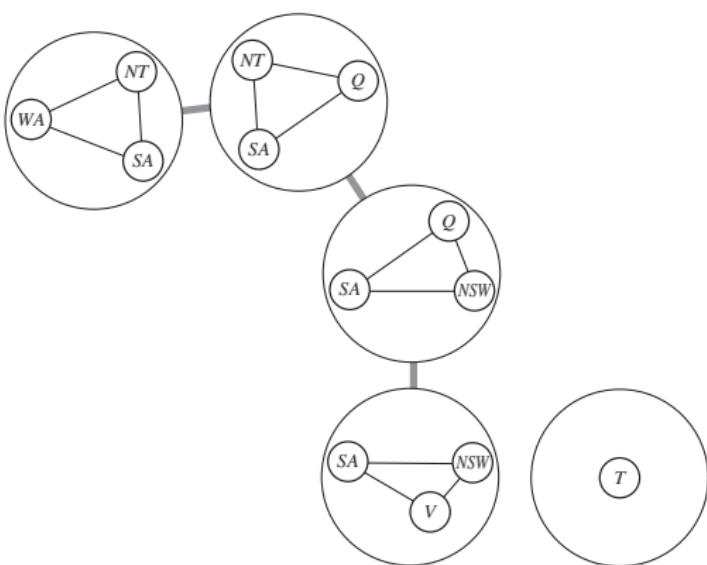
- Size of S is c : runtime $\mathcal{O}(d^c \cdot (n - c)d^2)$.
Explanation: We try each of the d^c combinations in S , and for each one we solve a tree problem of size $n - c$.
- For small c , the approach is very fast. In the worst case, c can become as large as $n - 2$.
- Finding the smallest S to obtain a tree-structure is NP-hard, but several efficient approximation algorithms are known.

⇒ check "Washington_CSP" for more illustration

Nearly tree-structured CSPs: Tree Decomposition (1)

Tree Decomposition: decomposing the CSP into a set of connected subproblems, connected by a “super-tree”.

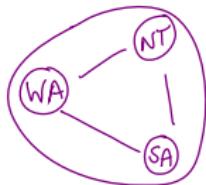
Requirements:



- Every variable X_i in the original problem appears in at least one subproblem.
- If two variables X_i and X_j are connected in the original problem, they must be connected in at least one subproblem.
- If a variable appears in two subproblems in the super-tree, it must appear in every subproblem along the path connecting those subproblems.

Nearly tree-structured CSPs: Tree Decomposition (2)

- We solve each subproblem independently.
 - ~~If any subproblem has no solution, we know the entire problem has no solution.~~
 - If we can solve all subproblems, we attempt to construct a global solution as follows:
 - We view each subproblem as a “super-variable” \hat{X}_i whose domain \hat{D}_i is the set of all solutions of the subproblem.
- Example (leftmost subproblem):**



$$\hat{D}_1 = \left\{ \begin{array}{l} \{ WA = red, SA = blue, NT = green \}, \\ \{ WA = red, SA = green, NT = blue \}, \\ \{ WA = green, SA = blue, NT = red \}, \\ \{ WA = green, SA = red, NT = blue \}, \\ \{ WA = blue, SA = red, NT = green \}, \\ \{ WA = blue, SA = green, NT = red \} \end{array} \right\}$$

Nearly tree-structured CSPs: Tree Decomposition (3)

- ② We solve the constraints connecting the subproblems, using the efficient algorithm for trees. The constraints between subproblems simply insist that the solutions of the subproblems agree on their shared variables.

Example (leftmost solution):

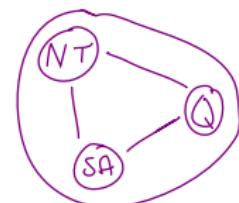
$$\{WA = \text{red}, SA = \text{blue}, NT = \text{green}\},$$

the only consistent solution for the next subproblem is

$$\{SA = \text{blue}, NT = \text{green}, Q = \text{red}\}$$

as can be verified using

$$\hat{D}_2 = \left\{ \begin{array}{l} \{SA = \text{red}, NT = \text{blue}, Q = \text{green}\}, \\ \{SA = \text{red}, NT = \text{green}, Q = \text{blue}\}, \\ \{SA = \text{green}, NT = \text{blue}, Q = \text{red}\}, \\ \{SA = \text{green}, NT = \text{red}, Q = \text{blue}\}, \\ \{SA = \text{blue}, NT = \text{red}, Q = \text{green}\}, \\ \{SA = \text{blue}, NT = \text{green}, Q = \text{red}\}, \end{array} \right\}$$



Summary

- states defined by values of a fixed set of variables
- goal test defined by constraints on variables values

- **Constraint satisfaction problems** (CSPs) require finding evaluations of variables within their domains to satisfy a set of constraints.
- **Backtracking search**, a form of depth-first search, is commonly used for solving CSPs. Inference can be interwoven with search.
with legal variable assigned per node
- * The **minimum-remaining-values** and **degree** heuristics are domain-independent methods to choose the next variable in backtracking search.
- * The **least-constraining-value** heuristics helps in selecting the first value for a variable.
- The complexity of solving CSPs strongly depends on the structure of their constraint graph:
 - Tree-structured problems can be solved in linear time.
 - **Conditioning** and **tree decomposition** can partially transform a problem into a tree-structured one.